

CAF

C++ Actor Framework

Matthias Vallentin
UC Berkeley


2016 Berkeley C++ Summit

Hello World


Hello World

```
main() {  
  
}
```


Hello World




Hello World




Hello World



Hello World



Hello World



Hello World

```
main() {  
  
}
```

Hello World

```
#include <string>
#include <iostream>

#include "caf/all.hpp"

using namespace caf;
using namespace std;

behavior mirror() {
    return {
        [=](const string& str) {
            return string(str.rbegin(), str.rend());
        }
    };
}

void hello_world(event_based_actor* self, const actor& buddy) {
    self->request(buddy, chrono::seconds(10), "Hello World!").then(
        [=](const string& str) {
            cout << str << endl;
        }
    );
}

int main() {
    actor_system_config cfg;
    actor_system system{cfg};
    auto m = system.spawn(mirror);
    system.spawn(hello_world, m);
}
```

```
#include <string>
#include <iostream>
```

```
#include "caf/all.hpp"
```

```
using namespace caf;
using namespace std;
```

```
behavior mirror() {
    return {
        [=](const string& str) {
            return string(str.rbegin(), str.rend());
        }
    };
}
```

Function returning behavior is-a actor

```
void hello_world(event_based_actor* self, const actor& buddy) {
    self->request(buddy, chrono::seconds(10), "Hello World!").then(
        [=](const string& str) {
            cout << str << endl;
        }
    );
}
```

```
int main() {
    actor_system_config cfg;
    actor_system system{cfg};
    auto m = system.spawn(mirror);
    system.spawn(hello_world, m);
}
```

```
#include <string>
#include <iostream>
```

```
#include "caf/all.hpp"
```

```
using namespace caf;
using namespace std;
```

```
behavior mirror() {
    return {
        [=](const string& str) {
            return string(str.rbegin(), str.rend());
        }
    };
}
```

Function returning behavior is-a actor

Function with event_based_actor as first argument is-a actor

```
void hello_world(event_based_actor* self, const actor& buddy) {
    self->request(buddy, chrono::seconds(10), "Hello World!").then(
        [=](const string& str) {
            cout << str << endl;
        }
    );
}
```

```
int main() {
    actor_system_config cfg;
    actor_system system{cfg};
    auto m = system.spawn(mirror);
    system.spawn(hello_world, m);
}
```

```
#include <string>
#include <iostream>
```

```
#include "caf/all.hpp"
```

```
using namespace caf;
using namespace std;
```

```
behavior mirror() {
    return {
        [=](const string& str) {
            return string(str.rbegin(), str.rend());
        }
    };
}
```

Function returning behavior is-a actor

Function with event_based_actor as first argument is-a actor

```
void hello_world(event_based_actor* self, const actor& buddy) {
    self->request(buddy, chrono::seconds(10), "Hello World!").then(
        [=](const string& str) {
            cout << str << endl;
        }
    );
}
```

Actor implicitly terminates at end of scope

```
int main() {
    actor_system_config cfg;
    actor_system system{cfg};
    auto m = system.spawn(mirror);
    system.spawn(hello_world, m);
}
```

```
#include <string>
#include <iostream>

#include "caf/all.hpp"
```

```
using namespace caf;
using namespace std;
```

```
behavior mirror() {
    return {
        [=](const string& str) {
            return string(str.rbegin(), str.rend());
        }
    };
}
```

Function returning behavior is-a actor

Function with event_based_actor as first argument is-a actor

```
void hello_world(event_based_actor* self, const actor& buddy) {
    self->request(buddy, chrono::seconds(10), "Hello World!").then(
        [=](const string& str) {
            cout << str << endl;
        }
    );
}
```

Actor implicitly terminates at end of scope

```
int main() {
    actor_system_config cfg;
    actor_system system{cfg};
    auto m = system.spawn(mirror);
    system.spawn(hello_world, m);
}
```

Encapsulates state: worker threads, actors, etc.

```
#include <string>
#include <iostream>

#include "caf/all.hpp"
```

```
using namespace caf;
using namespace std;
```

```
behavior mirror() {
    return {
        [=](const string& str) {
            return string(str.rbegin(), str.rend());
        }
    };
}
```

Function returning `behavior` is-a actor

Function with `event_based_actor` as first argument is-a actor

```
void hello_world(event_based_actor* self, const actor& buddy) {
    self->request(buddy, chrono::seconds(10), "Hello World!").then(
        [=](const string& str) {
            cout << str << endl;
        }
    );
}
```

Actor implicitly terminates at end of scope

```
int main() {
    actor_system_config cfg;
    actor_system system{cfg};
    auto m = system.spawn(mirror);
    system.spawn(hello_world, m);
}
```

Encapsulates state: worker threads, actors, etc.


`~actor_system()` blocks until all actors have terminated / no more references

Outline

- Actor Model
- CAF
- Performance

Actor Model

- **Actor**: sequential unit of computation
- **Message**: n-ary typed tuple
- **Mailbox**: FIFO/queue of messages
- **Behavior**: function how to process next message



Actor Semantics

Actor Semantics

- Actors are **reactive**

Actor Semantics

- Actors are **reactive**
- In response to a message, an actor can do *any* of:

Actor Semantics

- Actors are **reactive**
- In response to a message, an actor can do *any* of:
 1. Creating (*spawn*) new actors

Actor Semantics

- Actors are **reactive**
- In response to a message, an actor can do *any* of:
 1. Creating (*spawn*) new actors
 2. Sending messages to other actors

Actor Semantics

- Actors are **reactive**
- In response to a message, an actor can do *any* of:
 1. Creating (*spawn*) new actors
 2. Sending messages to other actors
 3. Designating a behavior for the next message

Key Advantages

Key Advantages

- All actors execute **concurrently**

Key Advantages

- All actors execute **concurrently**
- **Asynchronous** (non-blocking) message passing

Key Advantages

- All actors execute **concurrently**
- **Asynchronous** (non-blocking) message passing
- **Network-transparent** communication

Key Advantages


- All actors execute **concurrently**
- **Asynchronous** (non-blocking) message passing
- **Network-transparent** communication
- Hierarchical **fault propagation**

Key Advantages


- All actors execute **concurrently**
- **Asynchronous** (non-blocking) message passing
- **Network-transparent** communication
- Hierarchical **fault propagation**
- Sound semantics: **no data races** by design

CAF


Application
Logic



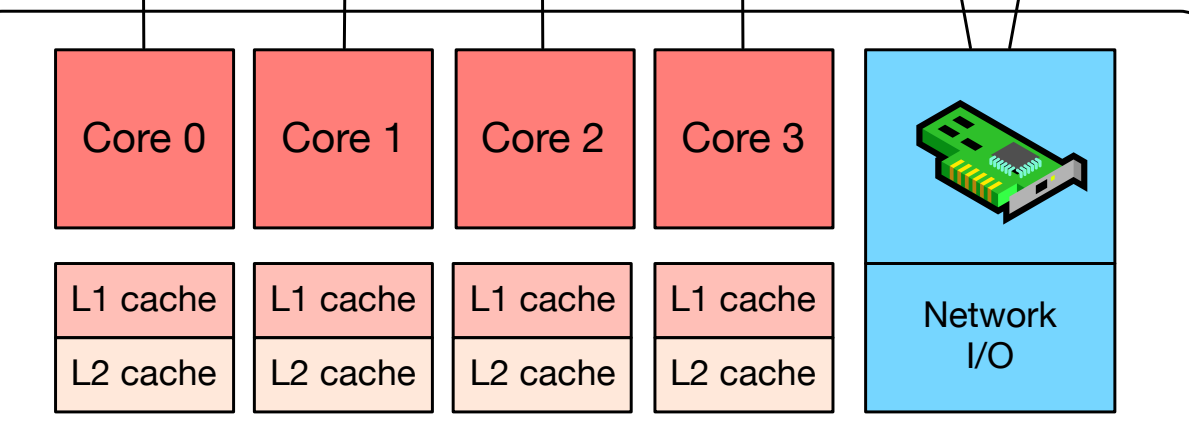
Actor
Runtime




Operating
System



Hardware

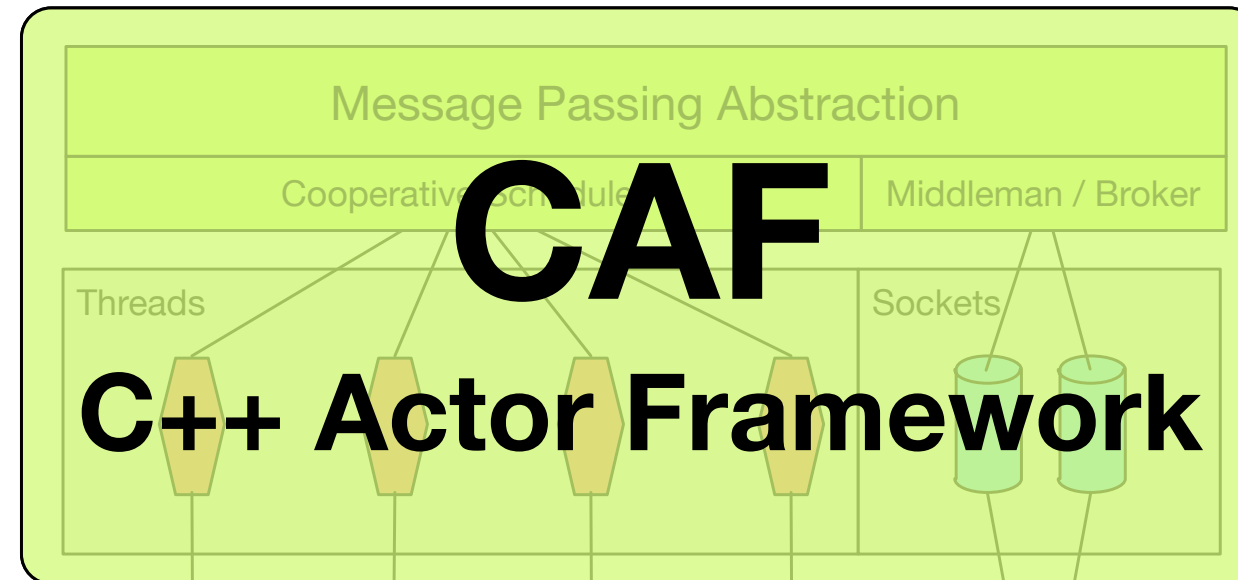


Application
Logic

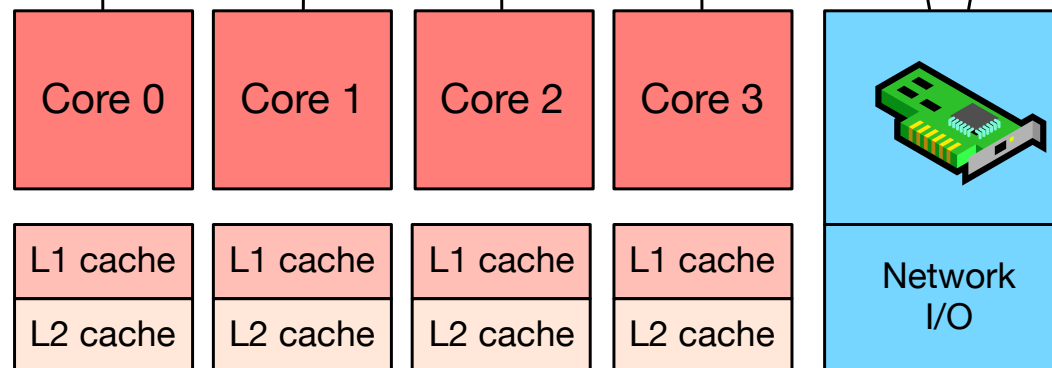


Actor
Runtime

Operating
System



Hardware



Scheduler

- Maps **N jobs** (= actors) to **M workers** (= threads)

- Maps **N jobs** (= actors) to **M workers** (= threads)
- Limitation: **cooperative multi-tasking** in user-space


- Maps **N jobs** (= actors) to **M workers** (= threads)
- Limitation: **cooperative multi-tasking** in user-space
- **Issue**: actors that block

- Maps **N jobs** (= actors) to **M workers** (= threads)
- Limitation: **cooperative multi-tasking** in user-space
- **Issue**: actors that block
 - Can lead to **starvation** and/or scheduling imbalance

- Maps **N jobs** (= actors) to **M workers** (= threads)
- Limitation: **cooperative multi-tasking** in user-space
- **Issue:** actors that block
 - Can lead to **starvation** and/or scheduling imbalance
 - Not well-suited for **I/O-heavy tasks**

- Maps **N jobs** (= actors) to **M workers** (= threads)
- Limitation: **cooperative multi-tasking** in user-space
- **Issue**: actors that block
 - Can lead to **starvation** and/or scheduling imbalance
 - Not well-suited for **I/O-heavy tasks**
 - Current solution: detach "uncooperative" actors into **separate thread**


Work Stealing*



*Robert D. Blumofe and Charles E. Leiserson. *Scheduling Multithreaded Computations by Work Stealing*. J. ACM, 46(5):720–748, September 1999.


Work Stealing*

- **Decentralized:** one job queue and worker thread per core




Work Stealing*

- **Decentralized**: one job queue and worker thread per core
- On empty queue, **steal** from other thread




Work Stealing*

- **Decentralized**: one job queue and worker thread per core
- On empty queue, **steal** from other thread




Work Stealing*

- **Decentralized**: one job queue and worker thread per core
- On empty queue, **steal** from other thread
- Efficient if stealing is a rare event



Work Stealing*


- **Decentralized**: one job queue and worker thread per core
- On empty queue, **steal** from other thread
- Efficient if stealing is a rare event
- Implementation: deque with two spinlocks



*Robert D. Blumofe and Charles E. Leiserson. *Scheduling Multithreaded Computations by Work Stealing*. J. ACM, 46(5):720–748, September 1999.

Work Stealing*


- **Decentralized**: one job queue and worker thread per core
- On empty queue, **steal** from other thread
- Efficient if stealing is a rare event
- Implementation: deque with two spinlocks
- **Graceful downscaling** of polling
 - **100** (0) \rightarrow **500** (10 μ s) \rightarrow ∞ (10ms)



Implementation


```
template <class Worker>
resumable* dequeue(Worker* self) {
    auto& strategies = self->data().strategies;
    resumable* job = nullptr;
    for (auto& strat : strategies) {
        for (size_t i = 0; i < strat.attempts; i += strat.step_size) {
            // try to grab a job from the front of the queue
            job = self->data().queue.take_head();
            // if we have still jobs, we're good to go
            if (job)
                return job;
            // try to steal every X poll attempts
            if ((i % strat.steal_interval) == 0) {
                if (job = try_steal(self))
                    return job;
            }
            if (strat.sleep_duration.count() > 0)
                std::this_thread::sleep_for(strat.sleep_duration);
        }
    }
    // unreachable, because the last strategy loops
    // until a job has been dequeued
    return nullptr;
}
```

Work Sharing




Work Sharing

- **Centralized:** one shared global queue




Work Sharing

- **Centralized**: one shared global queue
- Synchronization: **mutex & CV**




Work Sharing

- **Centralized**: one shared global queue
- Synchronization: **mutex & CV**
- **No polling**




Work Sharing

- **Centralized**: one shared global queue
- Synchronization: **mutex & CV**
- **No polling**
 - less CPU usage
 - lower throughput



Work Sharing

- **Centralized**: one shared global queue
- Synchronization: **mutex & CV**
- **No polling**
 - less CPU usage
 - lower throughput
- Good **for low-power devices**
 - Embedded / IoT



Copy-On-Write

- **Message** = atomic, intrusive
reference-counted tuple

- **Message** = atomic, intrusive reference-counted tuple
- **Immutable access** permitted

- **Message** = atomic, intrusive reference-counted tuple
- **Immutable access** permitted
- **Mutable access** with reference count > 1 invokes copy constructor

- **Message** = atomic, intrusive reference-counted tuple
- **Immutable access** permitted
- **Mutable access** with reference count > 1 invokes copy constructor
- **Constness deduced** from message handlers

- **Message** = atomic, intrusive reference-counted tuple
- **Immutable access** permitted
- **Mutable access** with reference count > 1 invokes copy constructor
- **Constness deduced** from message handlers

```
auto heavy = vector<char>(1024 * 1024);
auto msg = make_message(move(heavy));
for (auto& r : receivers)
    send(r, msg);
```

```
behavior reader() {
    return {
        [=](const vector<char>& buf) {
            f(buf);
        }
    };
}
```

```
behavior writer() {
    return {
        [=](vector<char>& buf) {
            f(buf);
        }
    };
}
```

- **Message** = atomic, intrusive reference-counted tuple
- **Immutable access** permitted
- **Mutable access** with reference count > 1 invokes copy constructor
- **Constness deduced** from message handlers

```
auto heavy = vector<char>(1024 * 1024);
auto msg = make_message(move(heavy));
for (auto& r : receivers)
    send(r, msg);
```

```
behavior reader() {
    return {
        [=](const vector<char>& buf) {
            f(buf);
```

const access enables efficient sharing of messages

```
behavior writer() {
    return {
        [=](vector<char>& buf) {
            f(buf);
        }
    };
}
```


- **Message** = atomic, intrusive reference-counted tuple
- **Immutable access** permitted
- **Mutable access** with reference count > 1 invokes copy constructor
- **Constness deduced** from message handlers

```
auto heavy = vector<char>(1024 * 1024);  
auto msg = make_message(move(heavy));  
for (auto& r : receivers)  
    send(r, msg);
```

```
behavior reader() {  
    return {  
        [=](const vector<char>& buf) {  
            f(buf);  
        }  
    }  
}
```

const access enables efficient sharing of messages

```
behavior writer() {  
    return {  
        [=](vector<char>& buf) {  
            f(buf);  
        }  
    }  
}
```

non-const access copies message contents *iff* ref count > 1

- **Message** = atomic, intrusive reference-counted tuple
- **Immutable access** permitted
- **Mutable access** with reference count > 1 invokes copy constructor
- **Constness deduced** from message handlers
- **No data races** by design

```
auto heavy = vector<char>(1024 * 1024);  
auto msg = make_message(move(heavy));  
for (auto& r : receivers)  
    send(r, msg);
```

```
behavior reader() {  
    return {  
        [=](const vector<char>& buf) {  
            f(buf);  
        }  
    };  
}
```

```
behavior writer() {  
    return {  
        [=](vector<char>& buf) {  
            f(buf);  
        }  
    };  
}
```

- **Message** = atomic, intrusive reference-counted tuple
- **Immutable access** permitted
- **Mutable access** with reference count > 1 invokes copy constructor
- **Constness deduced** from message handlers
- **No data races** by design
- **Value semantics**, no complex lifetime management

```
auto heavy = vector<char>(1024 * 1024);  
auto msg = make_message(move(heavy));  
for (auto& r : receivers)  
    send(r, msg);
```

```
behavior reader() {  
    return {  
        [=](const vector<char>& buf) {  
            f(buf);  
        }  
    };  
}
```

```
behavior writer() {  
    return {  
        [=](vector<char>& buf) {  
            f(buf);  
        }  
    };  
}
```

Type Safety

- CAF has **statically** and **dynamically typed** actors

- CAF has **statically** and **dynamically typed** actors
- **Dynamic**
 - Type-erased `caf::message` hides tuple types
 - Good for rapid prototyping

- CAF has **statically** and **dynamically typed** actors
- **Dynamic**
 - Type-erased `caf::message` hides tuple types
 - Good for rapid prototyping
- **Static**
 - Message protocol checked **at compile time**
 - **Type signature** verified at sender and receiver

Interface


```
// Atom: typed integer with semantics
using plus_atom = atom_constant<atom( "plus" )>;
using minus_atom = atom_constant<atom( "minus" )>;
using result_atom = atom_constant<atom( "result" )>;

// Actor type definition
using math_actor =
    typed_actor<
        replies_to<plus_atom, int, int>::with<result_atom, int>,
        replies_to<minus_atom, int, int>::with<result_atom, int>
    >;
```

Interface

```
// Atom: typed integer with semantics
using plus_atom = atom_constant<atom( "plus" )>;
using minus_atom = atom_constant<atom( "minus" )>;
using result_atom = atom_constant<atom( "result" )>;

// Actor type definition
using math_actor =
    typed_actor<
        replies_to<plus_atom, int, int>::with<result_atom, int>,
        replies_to<minus_atom, int, int>::with<result_atom, int>
    >;
```



Signature of **incoming** message

Interface

```
// Atom: typed integer with semantics  
using plus_atom = atom_constant<atom( "plus" )>;  
using minus_atom = atom_constant<atom( "minus" )>;  
using result_atom = atom_constant<atom( "result" )>;  
  
// Actor type definition  
using math_actor =  
    typed_actor<  
        replies_to<plus_atom, int, int>::with<result_atom, int>,  
        replies_to<minus_atom, int, int>::with<result_atom, int>  
    >;
```

Signature of **incoming** message

Signature of (optional) **response** message

Implementation

Implementation

Dynamic

```
behavior math_fun(event_based_actor* self) {  
  return {  
    [] (plus_atom, int a, int b) {  
      return make_tuple(result_atom::value, a + b);  
    },  
    [] (minus_atom, int a, int b) {  
      return make_tuple(result_atom::value, a - b);  
    }  
  };  
}
```

Implementation

Dynamic

```
behavior math_fun(event_based_actor* self) {  
  return {  
    [] (plus_atom, int a, int b) {  
      return make_tuple(result_atom::value, a + b);  
    },  
    [] (minus_atom, int a, int b) {  
      return make_tuple(result_atom::value, a - b);  
    }  
  };  
}
```

Static

```
math_actor::behavior_type typed_math_fun(math_actor::pointer self) {  
  return {  
    [] (plus_atom, int a, int b) {  
      return make_tuple(result_atom::value, a + b);  
    },  
    [] (minus_atom, int a, int b) {  
      return make_tuple(result_atom::value, a - b);  
    }  
  };  
}
```

Error Example

```
auto self = sys.spawn(...);  
math_actor m = self->typed_spawn(typed_math);  
self->request(m, seconds(1), plus_atom::value, 10, 20).then(  
    [](result_atom, float result) {  
  
    }  
);
```

Error Example

```
auto self = sys.spawn(...);  
math_actor m = self->typed_spawn(typed_math);  
self->request(m, seconds(1), plus_atom::value, 10, 20).then(  
    [](result_atom, float result) {  
    }  
);
```

Compiler complains about invalid response type

Class-Based API

Class-Based API

```
class math : public event_based_actor {  
public:  
    behavior make_behavior() override {  
        return {  
            [] (plus_atom, int a, int b) {  
                return make_tuple(result_atom::value, a + b);  
            },  
            [] (minus_atom, int a, int b) {  
                return make_tuple(result_atom::value, a - b);  
            }  
        };  
    }  
};
```

Dynamic

Class-Based API


Dynamic

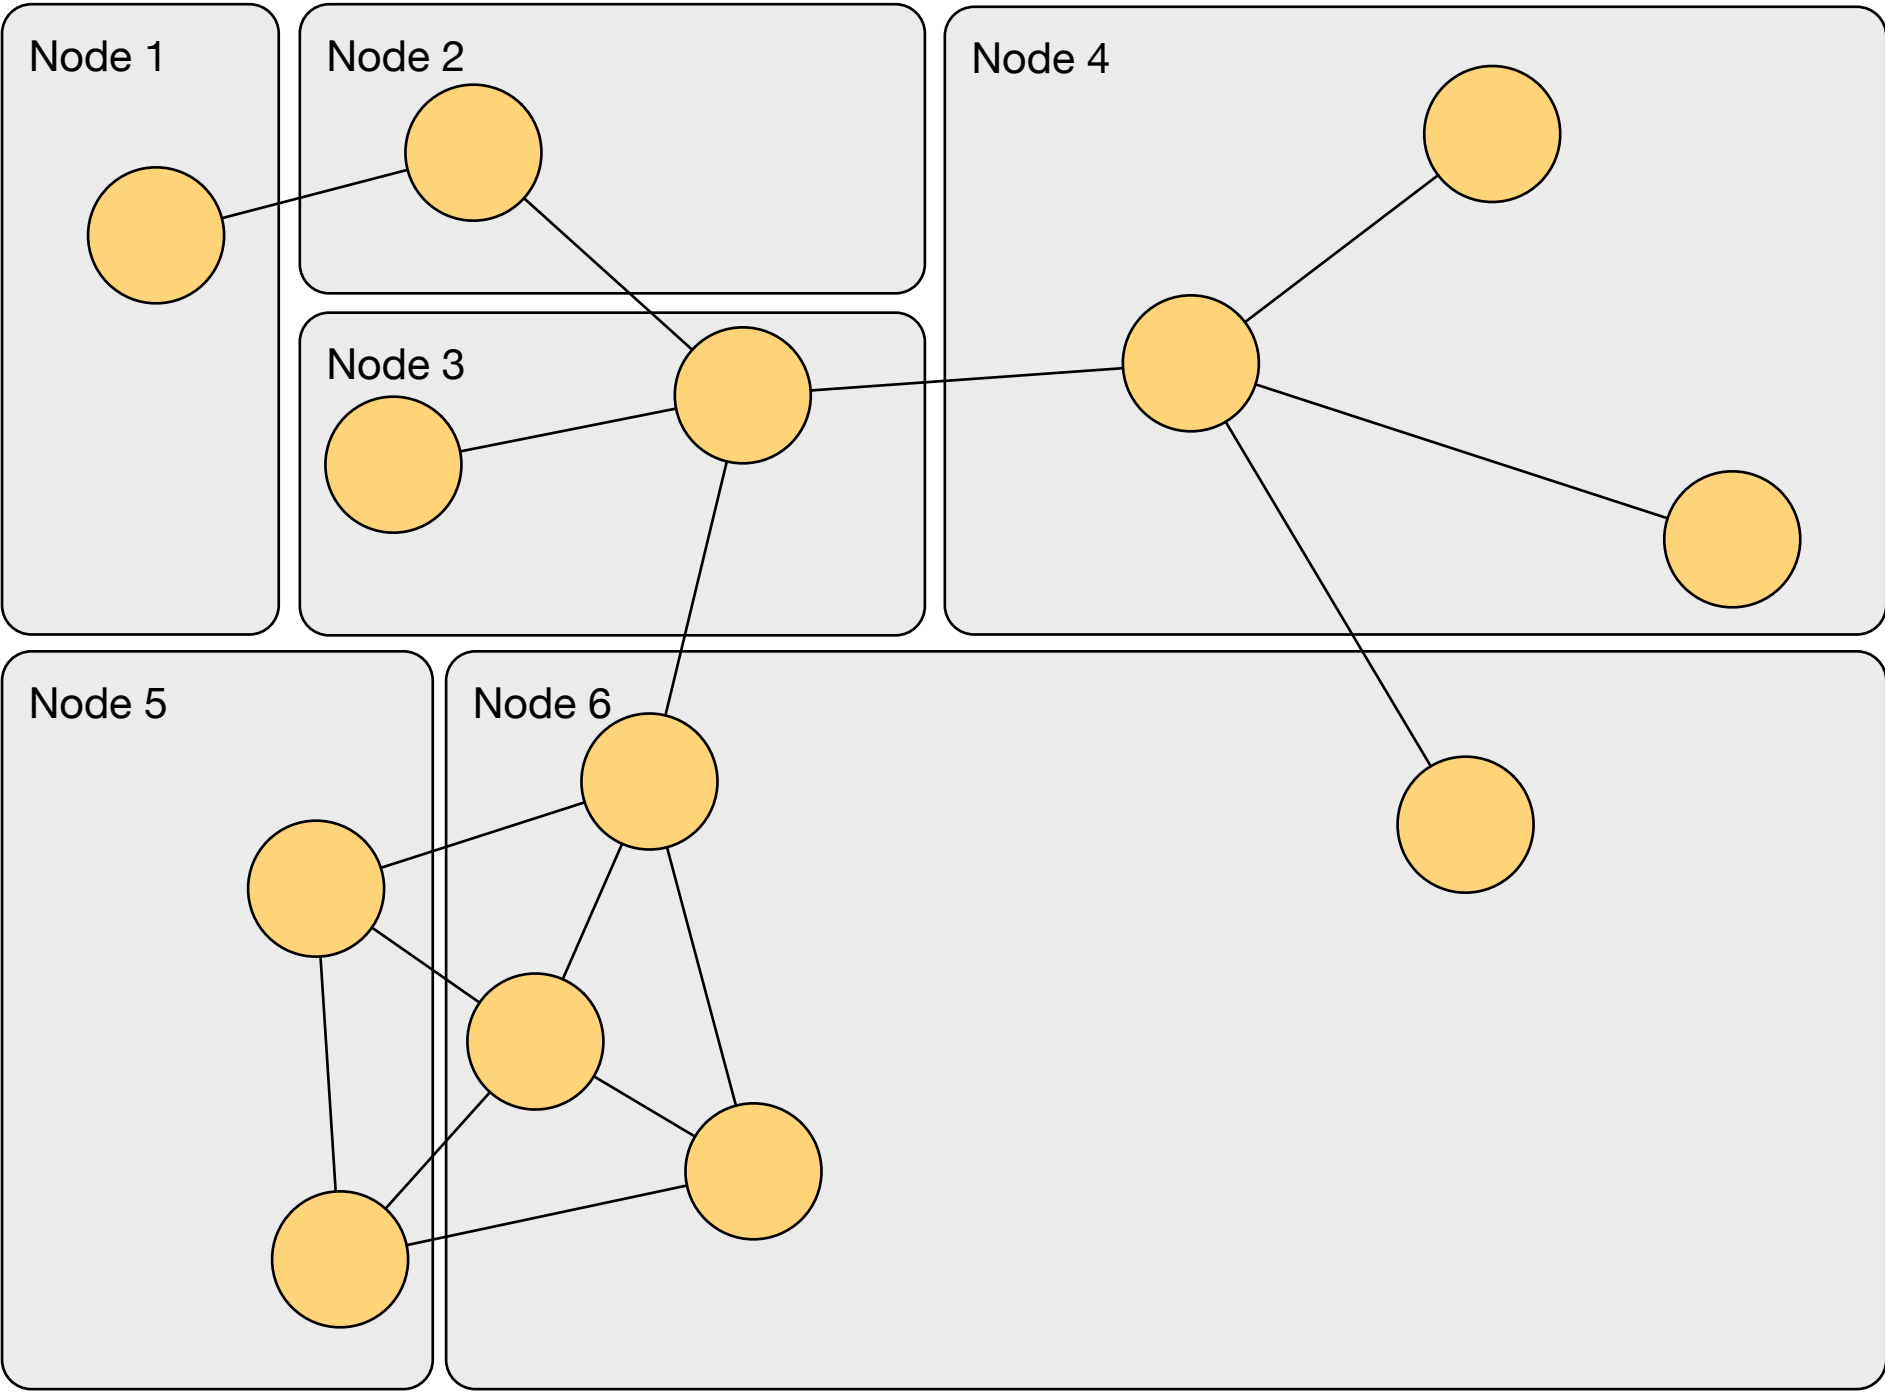
```
class math : public event_based_actor {  
public:  
    behavior make_behavior() override {  
        return {  
            [](plus_atom, int a, int b) {  
                return make_tuple(result_atom::value, a + b);  
            },  
            [](minus_atom, int a, int b) {  
                return make_tuple(result_atom::value, a - b);  
            }  
        };  
    }  
};
```

Static


```
class typed_math : public math_actor::base {  
public:  
    behavior_type make_behavior() override {  
        return {  
            [](plus_atom, int a, int b) {  
                return make_tuple(result_atom::value, a + b);  
            },  
            [](minus_atom, int a, int b) {  
                return make_tuple(result_atom::value, a - b);  
            }  
        };  
    }  
};
```

Network Transparency





Node 1



Network Transparency

Network Transparency

Separation of **application logic** from **deployment**

Network Transparency

Separation of **application logic** from **deployment**

- Significant **productivity gains**
 - Spend *more time* with **domain-specific code**
 - Spend *less time* with **network glue code**

Example

```
int main(int argc, char** argv) {  
    // Defaults.  
    auto host = "localhost"s;  
    auto port = uint16_t{42000};  
    auto server = false;  
    actor_system sys{...}; // Parse command line and setup actor system.  
    auto& middleman = sys.middleman();  
    actor a;  
    if (server) {  
        a = sys.spawn(math);  
        auto bound = middleman.publish(a, port);  
        if (bound == 0)  
            return 1;  
    } else {  
        auto r = middleman.remote_actor(host, port);  
        if (!r)  
            return 1;  
        a = *r;  
    }  
    // Interact with actor a  
}
```

Example

```
int main(int argc, char** argv) {  
    // Defaults.  
    auto host = "localhost"s;  
    auto port = uint16_t{42000};  
    auto server = false;  
    actor_system sys{...}; // Parse command line and setup actor system.  
    auto& middleman = sys.middleman();  
    actor a;  
    if (server) {  
        a = sys.spawn(math);  
        auto bound = middleman.publish(a, port);  
        if (bound == 0)  
            return 1;  
    } else {  
        auto r = middleman.remote_actor(host, port);  
        if (!r)  
            return 1;  
        a = *r;  
    }  
    // Interact with actor a  
}
```

Reference to CAF's network component.

Example

```
int main(int argc, char** argv) {  
    // Defaults.  
    auto host = "localhost"s;  
    auto port = uint16_t{42000};  
    auto server = false;  
    actor_system sys{...}; // Parse command line and setup actor system.  
    auto& middleman = sys.middleman();  
    actor a;  
    if (server) {  
        a = sys.spawn(math);  
        auto bound = middleman.publish(a, port);  
        if (bound == 0)  
            return 1;  
    } else {  
        auto r = middleman.remote_actor(host, port);  
        if (!r)  
            return 1;  
        a = *r;  
    }  
    // Interact with actor a  
}
```

Reference to CAF's network component.

Publish specific actor at a TCP port.
Returns bound port on success.

Example

```
int main(int argc, char** argv) {
```

```
    // Defaults.
```

```
    auto host = "localhost"s;
```

```
    auto port = uint16_t{42000};
```

```
    auto server = false;
```

```
    actor_system sys{...}; // Parse command line and setup actor system.
```

```
    auto& middleman = sys.middleman();
```

```
    actor a;
```

```
    if (server) {
```

```
        a = sys.spawn(math);
```

```
        auto bound = middleman.publish(a, port);
```

```
        if (bound == 0)
```

```
            return 1;
```

```
    } else {
```

```
        auto r = middleman.remote_actor(host, port);
```

```
        if (!r)
```

```
            return 1;
```

```
        a = *r;
```

```
    }
```

```
    // Interact with actor a
```

```
}
```

Reference to CAF's network component.

Publish specific actor at a TCP port.
Returns bound port on success.

Connect to published actor at TCP endpoint.
Returns `expected<actor>`.

Failures

Failure Management

Failure Management

Components fail regularly in large-scale systems

Failure Management

Components fail regularly in large-scale systems

- Actor model provides **monitors** and **links**

Failure Management

Components fail regularly in large-scale systems

- Actor model provides **monitors** and **links**
 - **Monitor**: subscribe to exit of actor (**unidirectional**)

Failure Management

Components fail regularly in large-scale systems

- Actor model provides **monitors** and **links**
 - **Monitor**: subscribe to exit of actor (**unidirectional**)
 - **Link**: bind own lifetime to other actor (**bidirectional**)

Failure Management


Components fail regularly in large-scale systems

- Actor model provides **monitors** and **links**
 - **Monitor**: subscribe to exit of actor (**unidirectional**)
 - **Link**: bind own lifetime to other actor (**bidirectional**)
- Enables **hierarchical propagation** of failures


Failure Management

Components fail regularly in large-scale systems


- Actor model provides **monitors** and **links**
 - **Monitor**: subscribe to exit of actor (**unidirectional**)
 - **Link**: bind own lifetime to other actor (**bidirectional**)
- Enables **hierarchical propagation** of failures
- Can define **local fault isolation** domains



Monitored actors: subscribe to notification if actor terminates



Monitored actors: subscribe to notification if actor terminates



Linked actors: either **all alive** or **collectively failing**

Monitor Example

```
behavior adder() {  
    return {  
        [](int x, int y) {  
            return x + y;  
        }  
    };  
}
```

```
auto self = sys.spawn<monitored>(adder);  
self->set_down_handler(  
    [](const down_msg& msg) {  
        cout << "actor DOWN: " << msg.reason << endl;  
    }  
);
```

Monitor Example

```
behavior adder() {  
    return {  
        [](int x, int y) {  
            return x + y;  
        }  
    };  
}
```

Spawn flag denotes monitoring.
Also possible later via **self->monitor(other);**

```
auto self = sys.spawn<monitored>(adder);  
self->set_down_handler(  
    [](const down_msg& msg) {  
        cout << "actor DOWN: " << msg.reason << endl;  
    }  
);
```

Link Example

```
behavior adder() {  
    return {  
        [](int x, int y) {  
            return x + y;  
        }  
    };  
}  
  
auto self = sys.spawn<linked>(adder);  
self->set_exit_handler(  
    [](const exit_msg& msg) {  
        cout << "actor EXIT: " << msg.reason << endl;  
    }  
);
```

Link Example

```
behavior adder() {  
    return {  
        [](int x, int y) {  
            return x + y;  
        }  
    };  
}
```

Spawn flag denotes linking.
Also possible later via **self->link_to(other);**

```
auto self = sys.spawn<linked>(adder);  
self->set_exit_handler(  
    [](const exit_msg& msg) {  
        cout << "actor EXIT: " << msg.reason << endl;  
    }  
);
```

Project Summary

Project Summary

- Lead: **Dominik Charousset** (HAW Hamburg)
 - Started CAF as Master's thesis
 - Active development as part of his Ph.D.

Project Summary

- Lead: **Dominik Charousset** (HAW Hamburg)
 - Started CAF as Master's thesis
 - Active development as part of his Ph.D.
- Dual-licensed: 3-clause **BSD** & **Boost**

Project Summary

- Lead: **Dominik Charousset** (HAW Hamburg)
 - Started CAF as Master's thesis
 - Active development as part of his Ph.D.
- Dual-licensed: 3-clause **BSD** & **Boost**
- Fast growing community (~1K stars on github, active ML)

Project Summary

- Lead: **Dominik Charousset** (HAW Hamburg)
 - Started CAF as Master's thesis
 - Active development as part of his Ph.D.
- Dual-licensed: 3-clause **BSD** & **Boost**
- Fast growing community (~1K stars on github, active ML)
- Presented CAF twice at C++Now
 - Feedback resulted in **type-safe actors**


Project Summary

- Lead: **Dominik Charousset** (HAW Hamburg)
 - Started CAF as Master's thesis
 - Active development as part of his Ph.D.
- Dual-licensed: 3-clause **BSD** & **Boost**
- Fast growing community (~1K stars on github, active ML)
- Presented CAF twice at C++Now
 - Feedback resulted in **type-safe actors**
- Production-grade code: extensive unit tests, comprehensive CI

Evaluation


<https://github.com/actor-framework/benchmarks>

Setup #1




Setup #1

- 100 rings of 100 actors each




Setup #1

- 100 rings of 100 actors each
- Actors forward single token 1K times, then terminate




Setup #1

- 100 rings of 100 actors each
- Actors forward single token 1K times, then terminate
- 4 re-creations per ring




Setup #1




- 100 rings of 100 actors each
- Actors forward single token 1K times, then terminate
- 4 re-creations per ring
- One actor per ring performs *prime factorization*

Setup #1




- 100 rings of 100 actors each
- Actors forward single token 1K times, then terminate
- 4 re-creations per ring
- One actor per ring performs *prime factorization*
- Resulting workload: high message & CPU pressure

Setup #1




- 100 rings of 100 actors each
- Actors forward single token 1K times, then terminate
- 4 re-creations per ring
- One actor per ring performs *prime factorization*
- Resulting workload: high message & CPU pressure
- Ideal: 2 x cores \Rightarrow 0.5 x runtime


Performance




(normalized)




(normalized)



Memory Overhead




Setup #2




Setup #2

- Compute images of Mandelbrot set




Setup #2

- Compute images of Mandelbrot set
- Divide & conquer algorithm




Setup #2

- Compute images of Mandelbrot set
- Divide & conquer algorithm
- Compare against OpenMPI (via Boost.MPI)
 - Only message passing layers differ




Setup #2


- Compute images of Mandelbrot set
- Divide & conquer algorithm
- Compare against OpenMPI (via Boost.MPI)
 - Only message passing layers differ
- 16-node cluster: quad-core Intel i7 3.4 GHz



CAF vs. OpenMPI



Scalability



Summary

- Actor model is a natural fit for today's systems
- CAF offers an efficient C++ runtime
 - High-level message passing abstraction
 - Type-safe messaging APIs *at compile time*
 - Network-transparent communication
 - Well-defined failure semantics

Questions?

<http://actor-framework.org>

<https://github.com/actor-framework>

Backup Slides

Sending Messages

- Asynchronous fire-and-forget

```
self->send(other, x, xs...);
```

- Request-response with one-shot handler

```
self->request(other, timeout, x, xs...).then(  
    [](T response) {  
    }  
);
```

- Transparent forwarding of message forwarding

```
self->delegate(other, x, xs...);
```

A Simple Actor

```
behavior adder(event_based_actor* self) {  
    return {  
        [] (int x, int y) {  
            return x + y;  
        }  
    };  
}
```

A Simple Actor

```
behavior adder(event_based_actor* self) {  
    return {  
        [] (int x, int y) {  
            return x + y;  
        }  
    };  
}
```

A non-void return value sends a message
back to the original sender

A Simple Actor

Optional reference to the running actor, e.g.,
to send messages in response handlers.

```
behavior adder(event_based_actor* self) {  
    return {  
        [] (int x, int y) {  
            return x + y;  
        }  
    };  
}
```

A non-void return value sends a message
back to the original sender

A Simple Actor

An actor is typically implemented as function

Optional reference to the running actor, e.g.,
to send messages in response handlers.

```
behavior adder(event_based_actor* self) {  
    return {  
        [] (int x, int y) {  
            return x + y;  
        }  
    };  
}
```

A non-void return value sends a message
back to the original sender

A Simple Actor

An actor is typically implemented as function

Optional reference to the running actor, e.g., to send messages in response handlers.

```
behavior adder(event_based_actor* self) {  
    return {  
        [] (int x, int y) {  
            return x + y;  
        }  
    };  
}
```

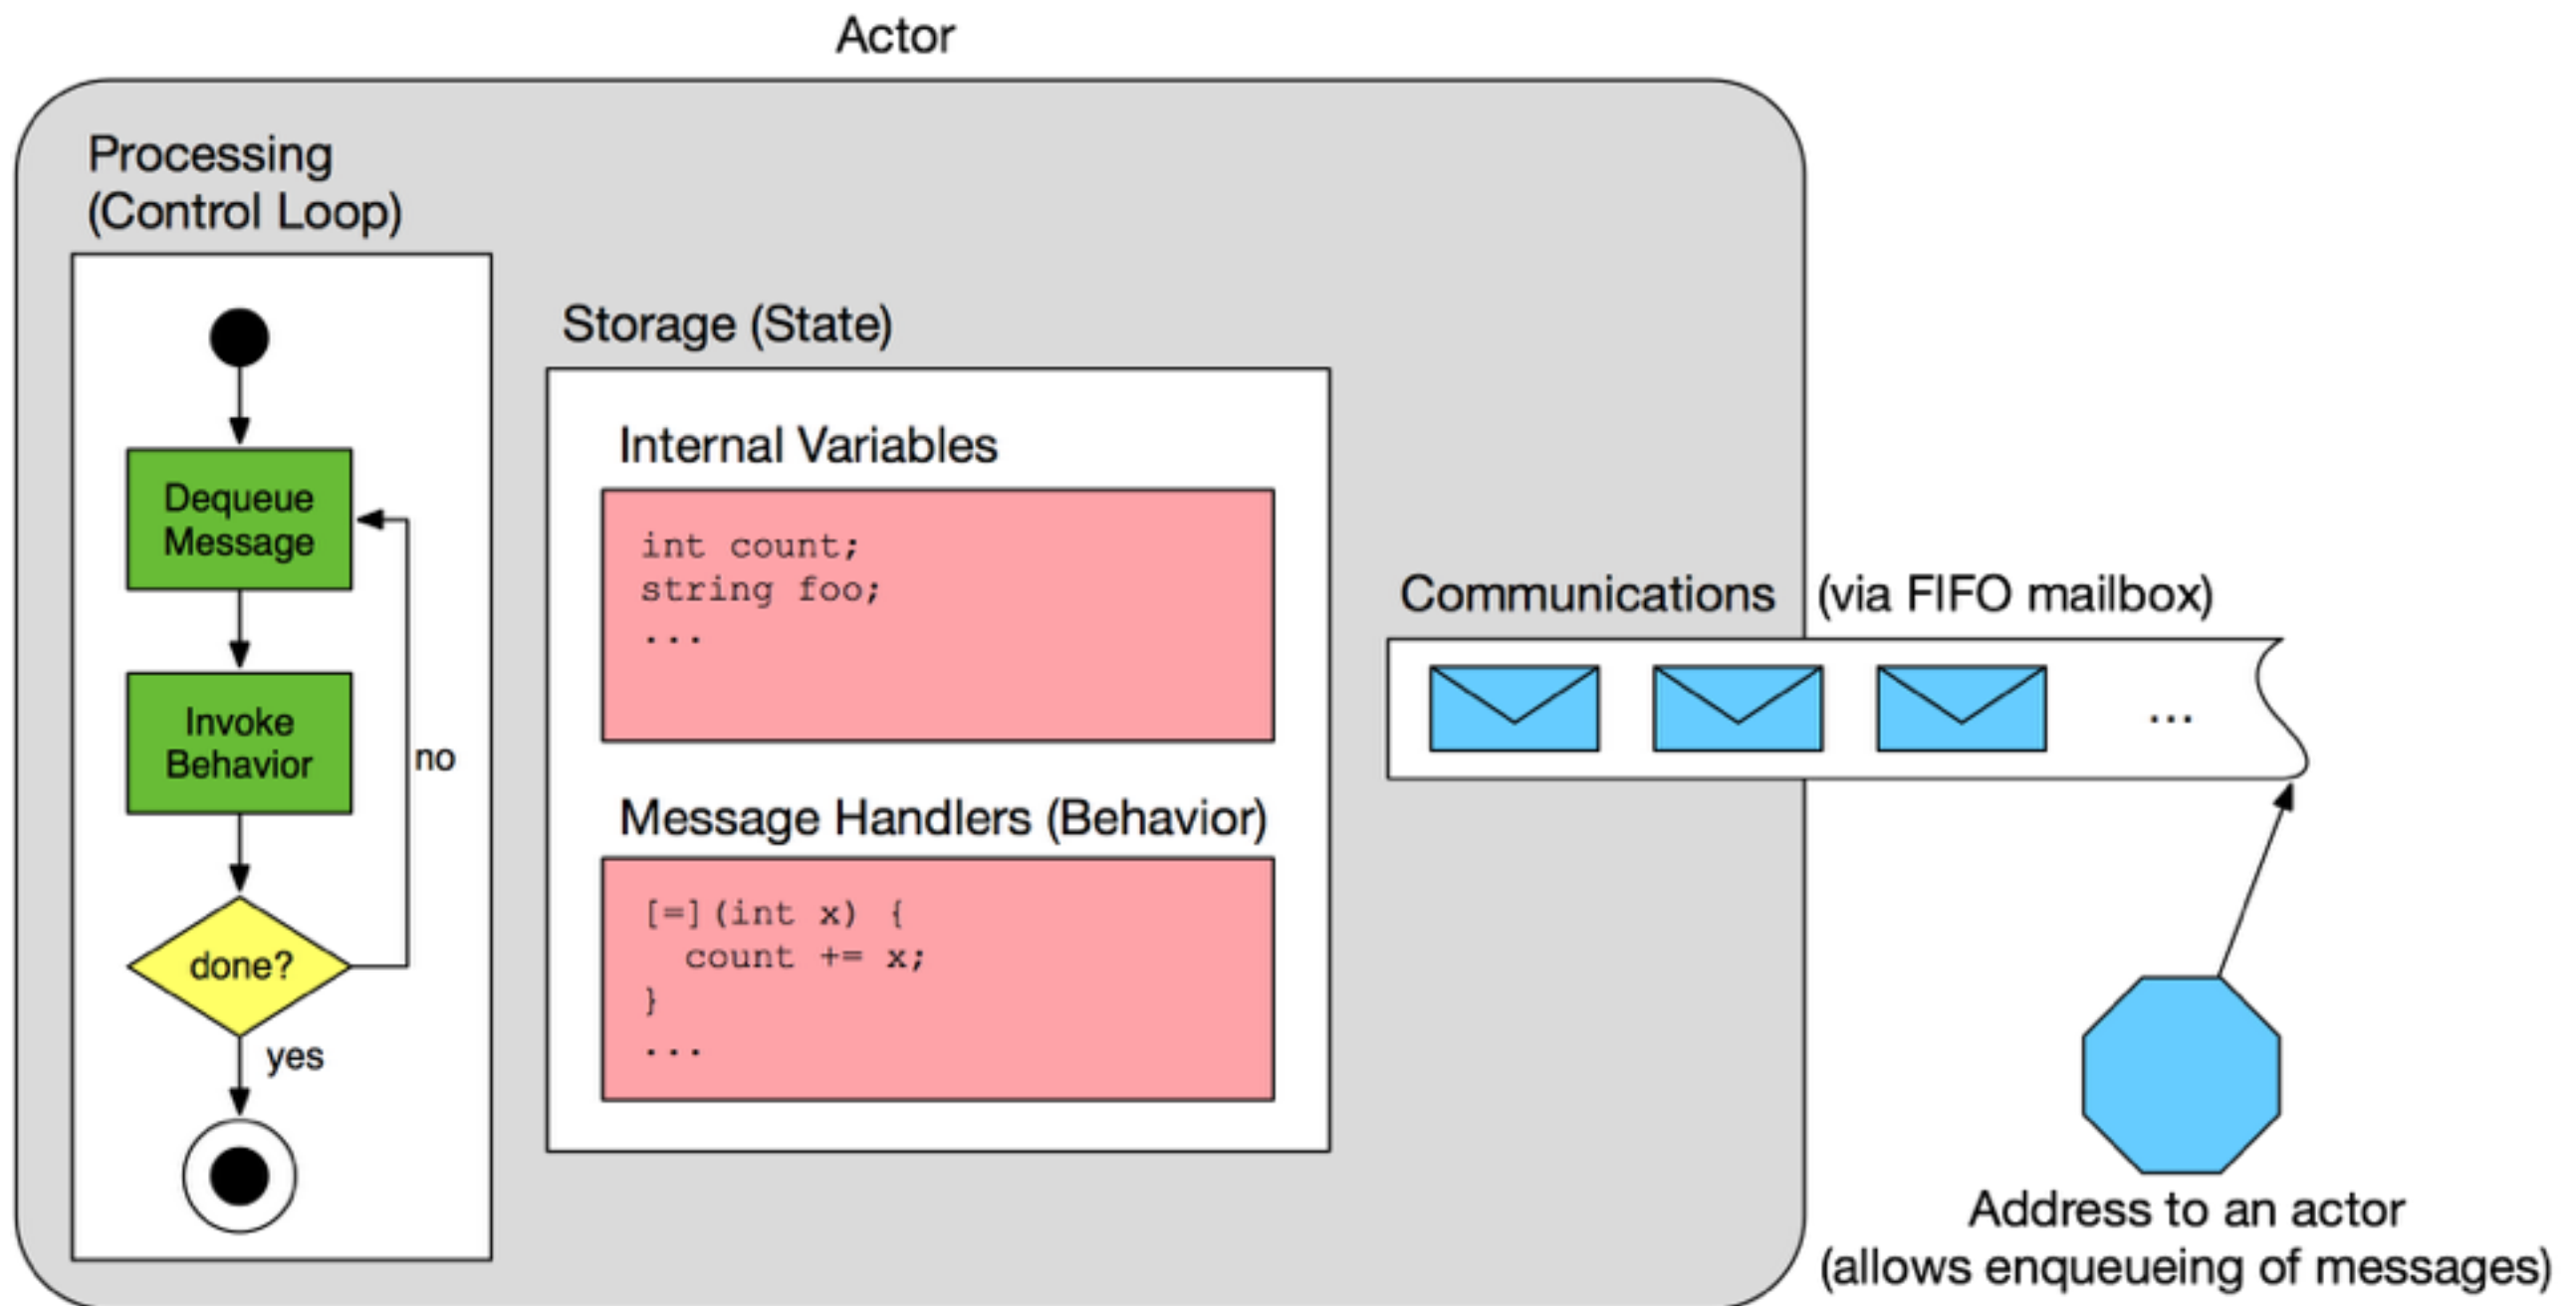
A list of lambda functions determines the message types an actor processes.

A non-void return value sends a message back to the original sender

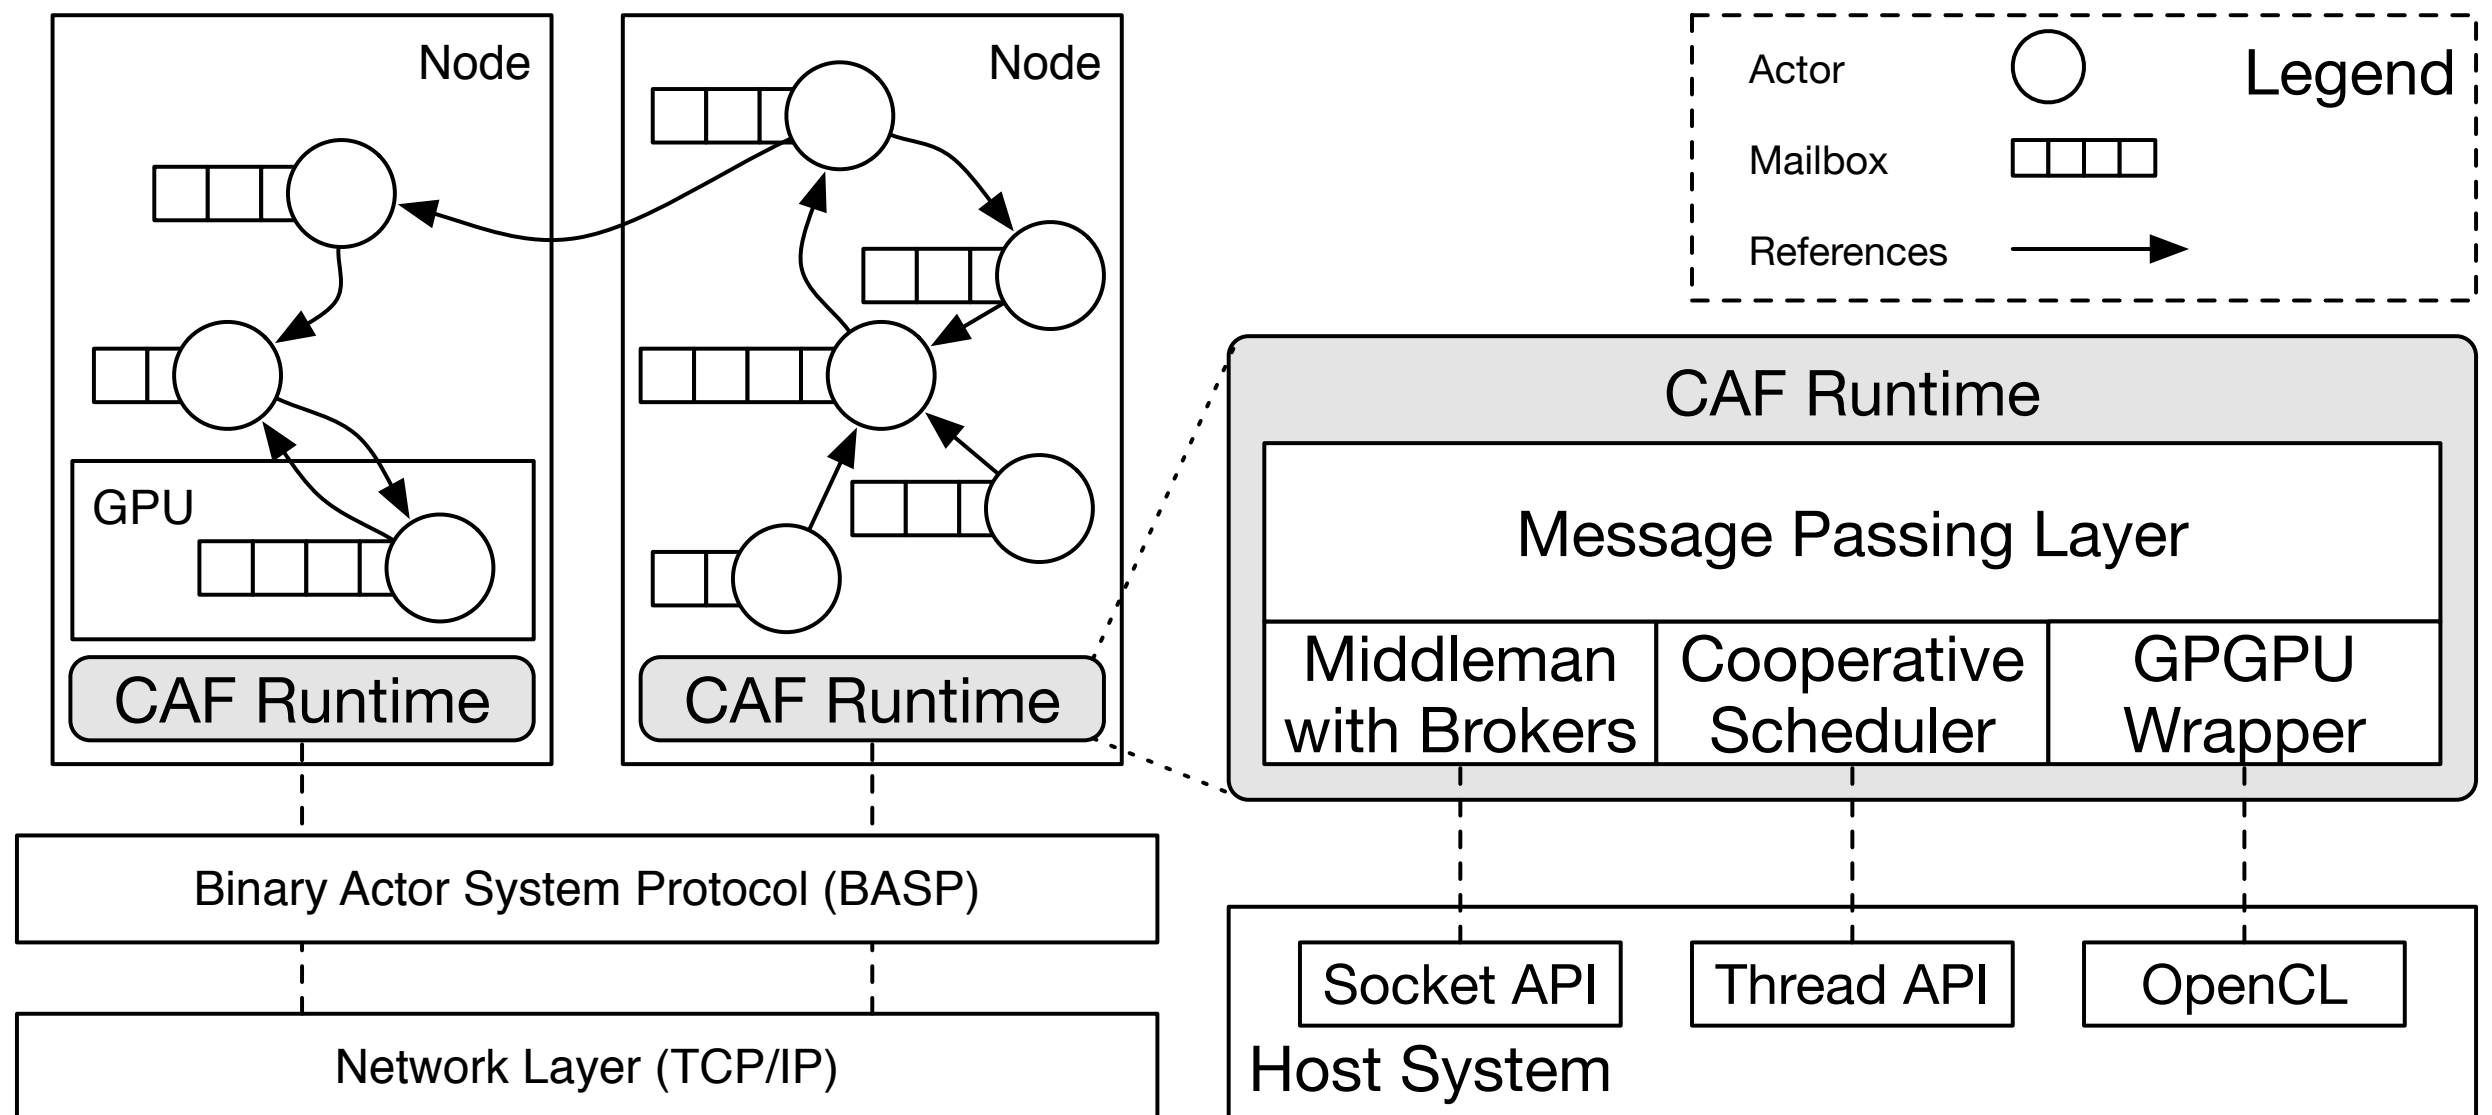
Actors as Function Objects

```
actor a = sys.spawn(adder);  
auto f = make_function_view(a);  
cout << "f(1, 2) = "  
      << to_string(f(1, 2))  
      << "\n";
```



Anatomy of an Actor




Architecture




Actor Creation




Actor Creation

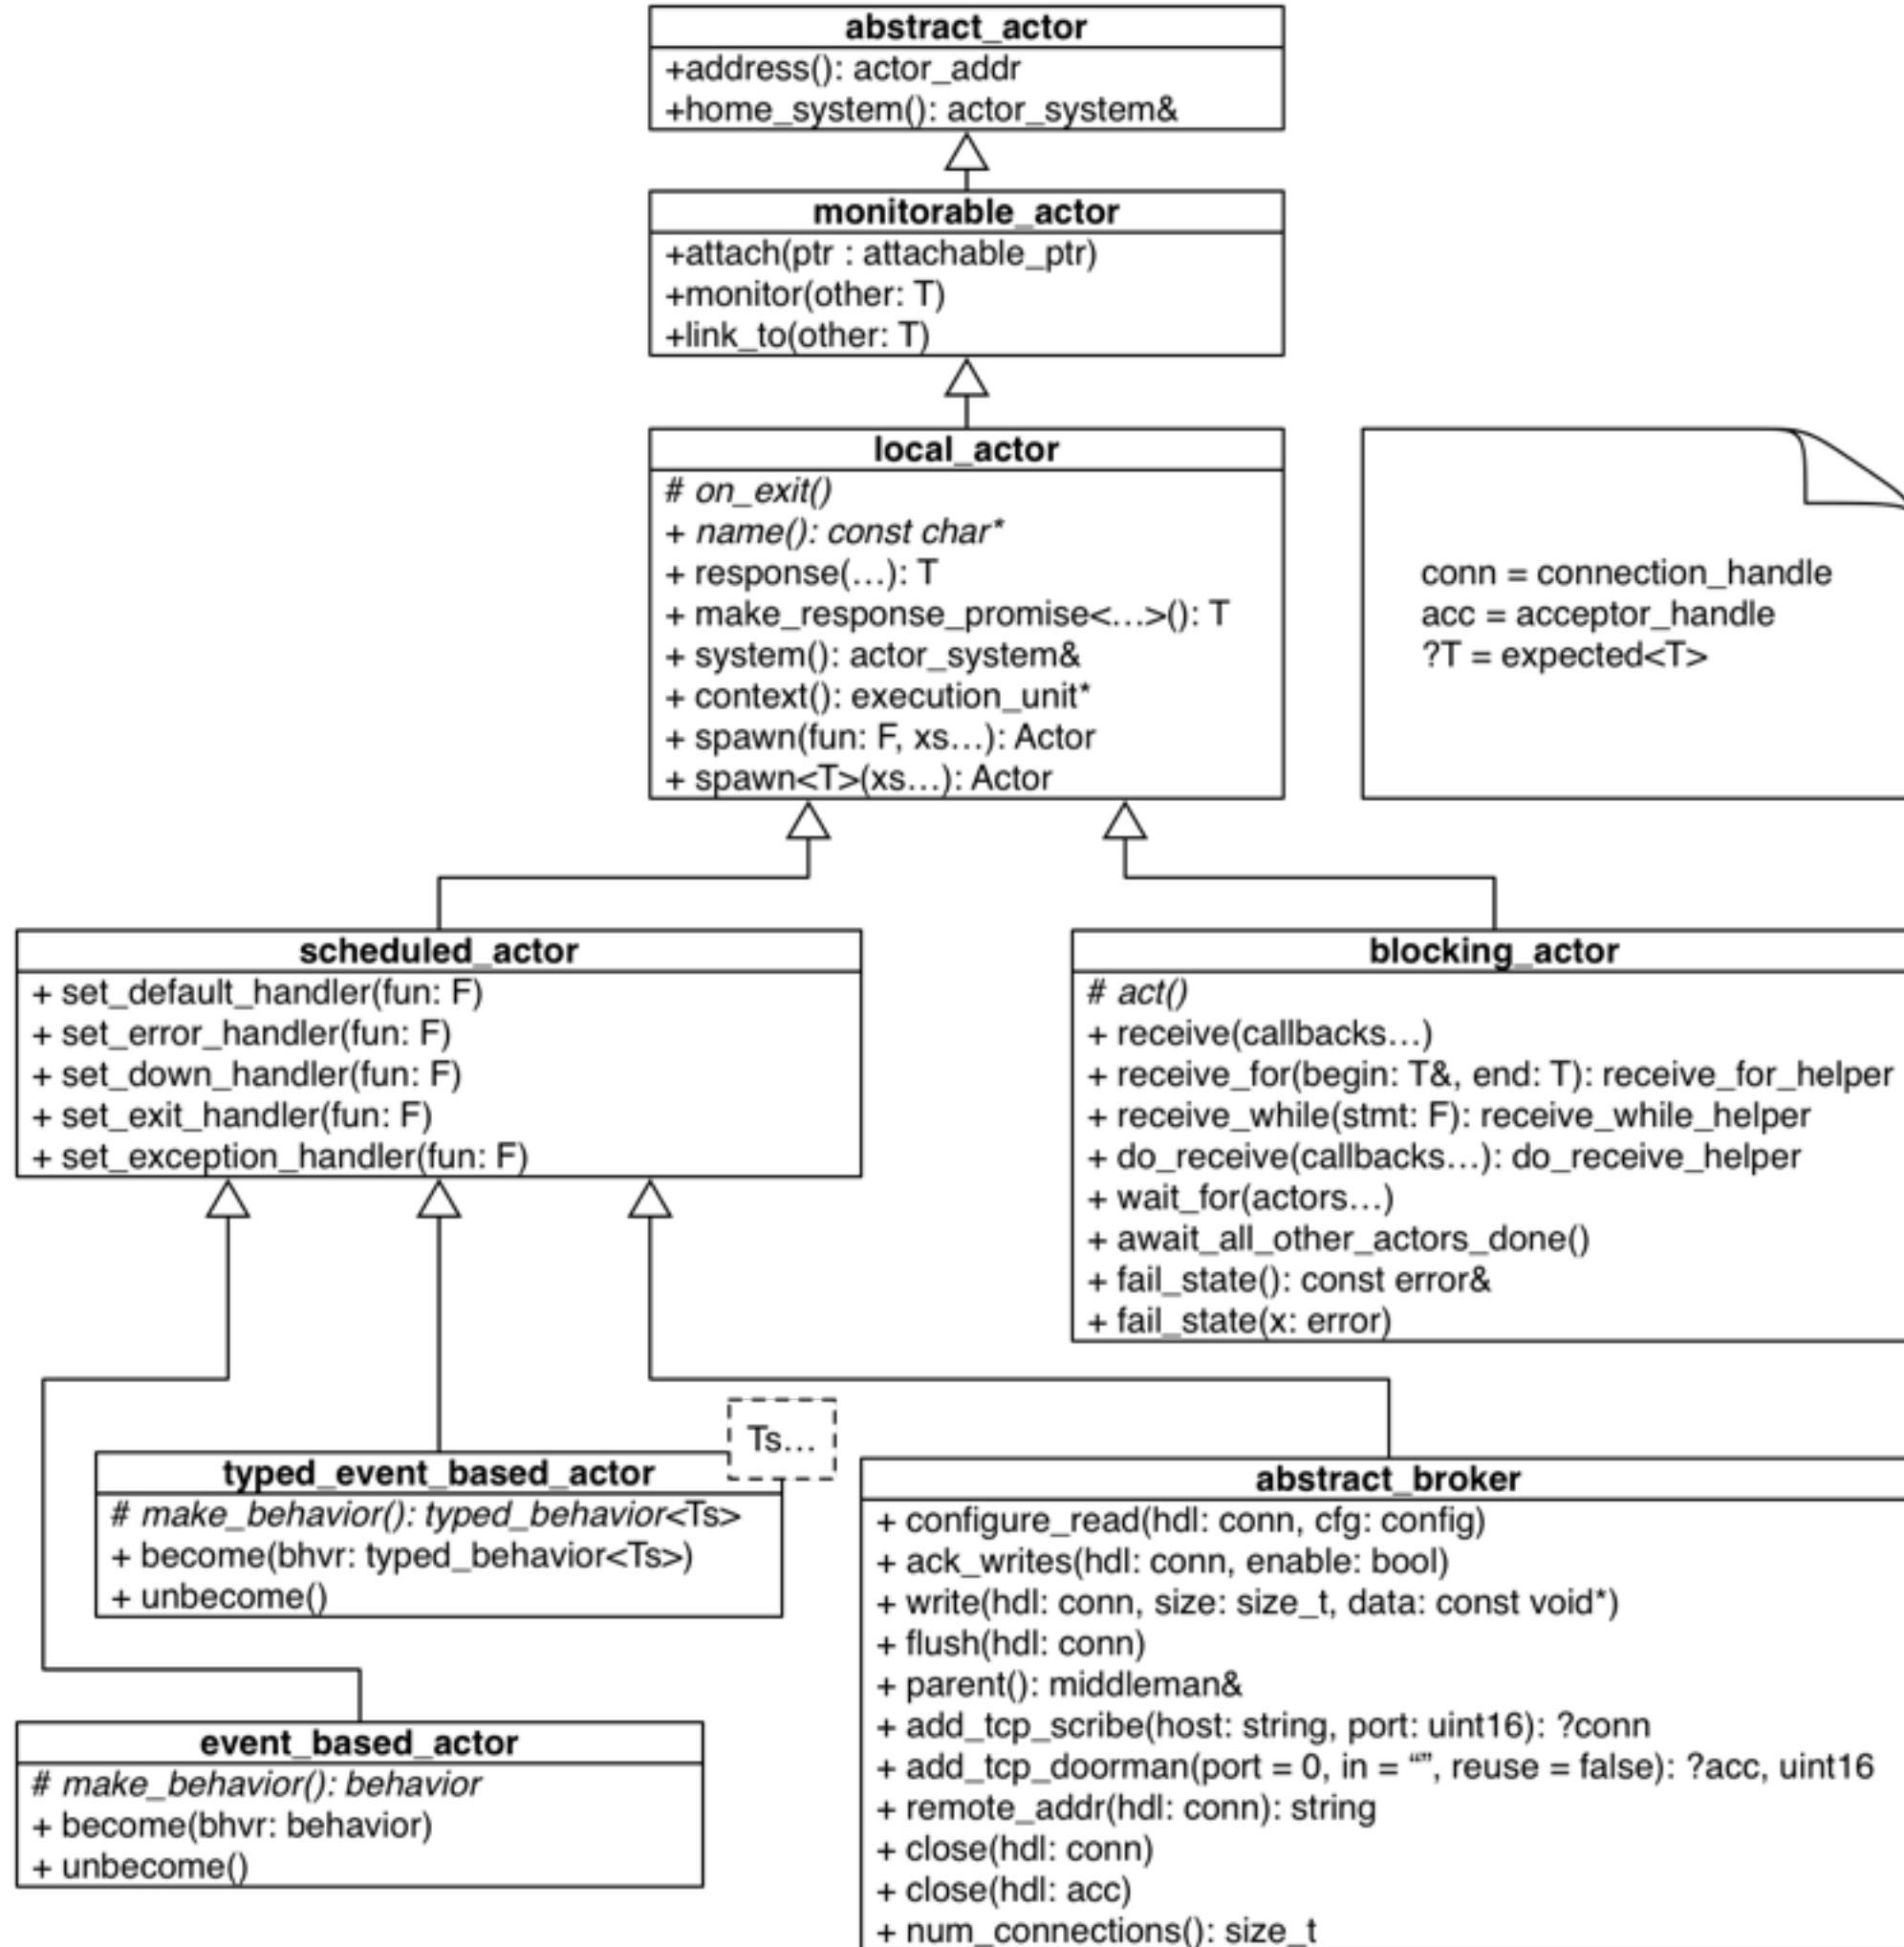


Mailbox - CPU



Mailbox - Memory





Native Execution	Garbage Collection	Pattern Matching	Copy-On-Write Messaging	Failure Propagation	Dynamic Behaviors	Compile-Time Type Checking	Run-Time Type Checking	Exchangeable Scheduler	Network Actors	GPU Actors	Backend
------------------	--------------------	------------------	-------------------------	---------------------	-------------------	----------------------------	------------------------	------------------------	----------------	------------	---------

Erlang [13]	✗	✓	✓	✗	✓	✓	✗	✓	✓	✓	✗	BEAM
Elixir [70]	✗	✓	✓	✗	✓	✓	✗	✓	✗	✓	✗	BEAM
Akka/Scala [7]	✗	✓	✓	✗	✓	✓	✓	✓	✓	✓	✗	JVM
SALSA Lite [62]	✗	✓	✗	✗	✗	✗	✓	✓	✗	● [†]	✗	JVM
Actor Foundry [2]	✗	✓	✓	✗	✗	✓	✗	✓	✗	✓	✗	JVM
Pulsar [151]	✗	✓	✓	✗	✓	✓	✗	✓	✗	✓	✗	JVM
Pony [47]	✓	✓	✓	✗	✗	✗	✓	✓	✗	✓	✗	LLVM
Charm++ [109]	✓	● [*]	✗	✗	✗	✗	✓	✓	✗	✓	✗	C++
Theron [182]	✓	● [*]	✗	✗	✗	✓	✗	✓	✗	✓	✗	C++
libprocess [124]	✓	● [*]	✗	✗	✓	✗	✓	✗	✗	✓	✗	C++
CAF [44]	✓	● [*]	✓	✓	✓	✓	✓	✓	✓	✓	✓	C++

* Via reference counting, as opposed to tracing garbage collection.

† Only in SALSA, not SALSA Lite.