

HPX

A C++ Standard Library for Parallelism and Concurrency

Make C++ Parallelism Independent of any External Solutions by 2020

Companions Left Behind

- Parallelism in C++ today depends on external solutions
 - Pragma based
 - OpenMP
 - OpenACC
 - External Libraries (mostly 'C' libraries)
 - CUDA
 - MPI
 - TBB
- Insufficient integration with
 - C++ type system and
 - C++ standard libraries
- Insufficient integration between those solutions

Types of Parallelism

- Data parallelism
 - Iterative fork/join
 - `#pragma omp parallel for`, very limited, not integrated with type system
 - Accelerators, coprocessors
 - OpenACC, various library solutions
- Task based asynchronous and continuation style parallelism
 - Fork/join of heterogeneous tasks
 - OpenMP tasks, difficult to compose, difficult to nest
 - Task-based parallelism
 - OpenMP tasks, very limited as they can't exit current scope
 - TBB, weakly aligned with C++ Standard library
 - Asynchronous continuation style
 - ???

Types of Parallelism

- Instruction Level Parallelism
 - Vectorization, SIMD instructions
 - `#pragma simd`
- Distributed computing
 - MPI, mostly data transfer, coarse grain coordination (barriers)
 - Difficult to use in conjunction with any of the above

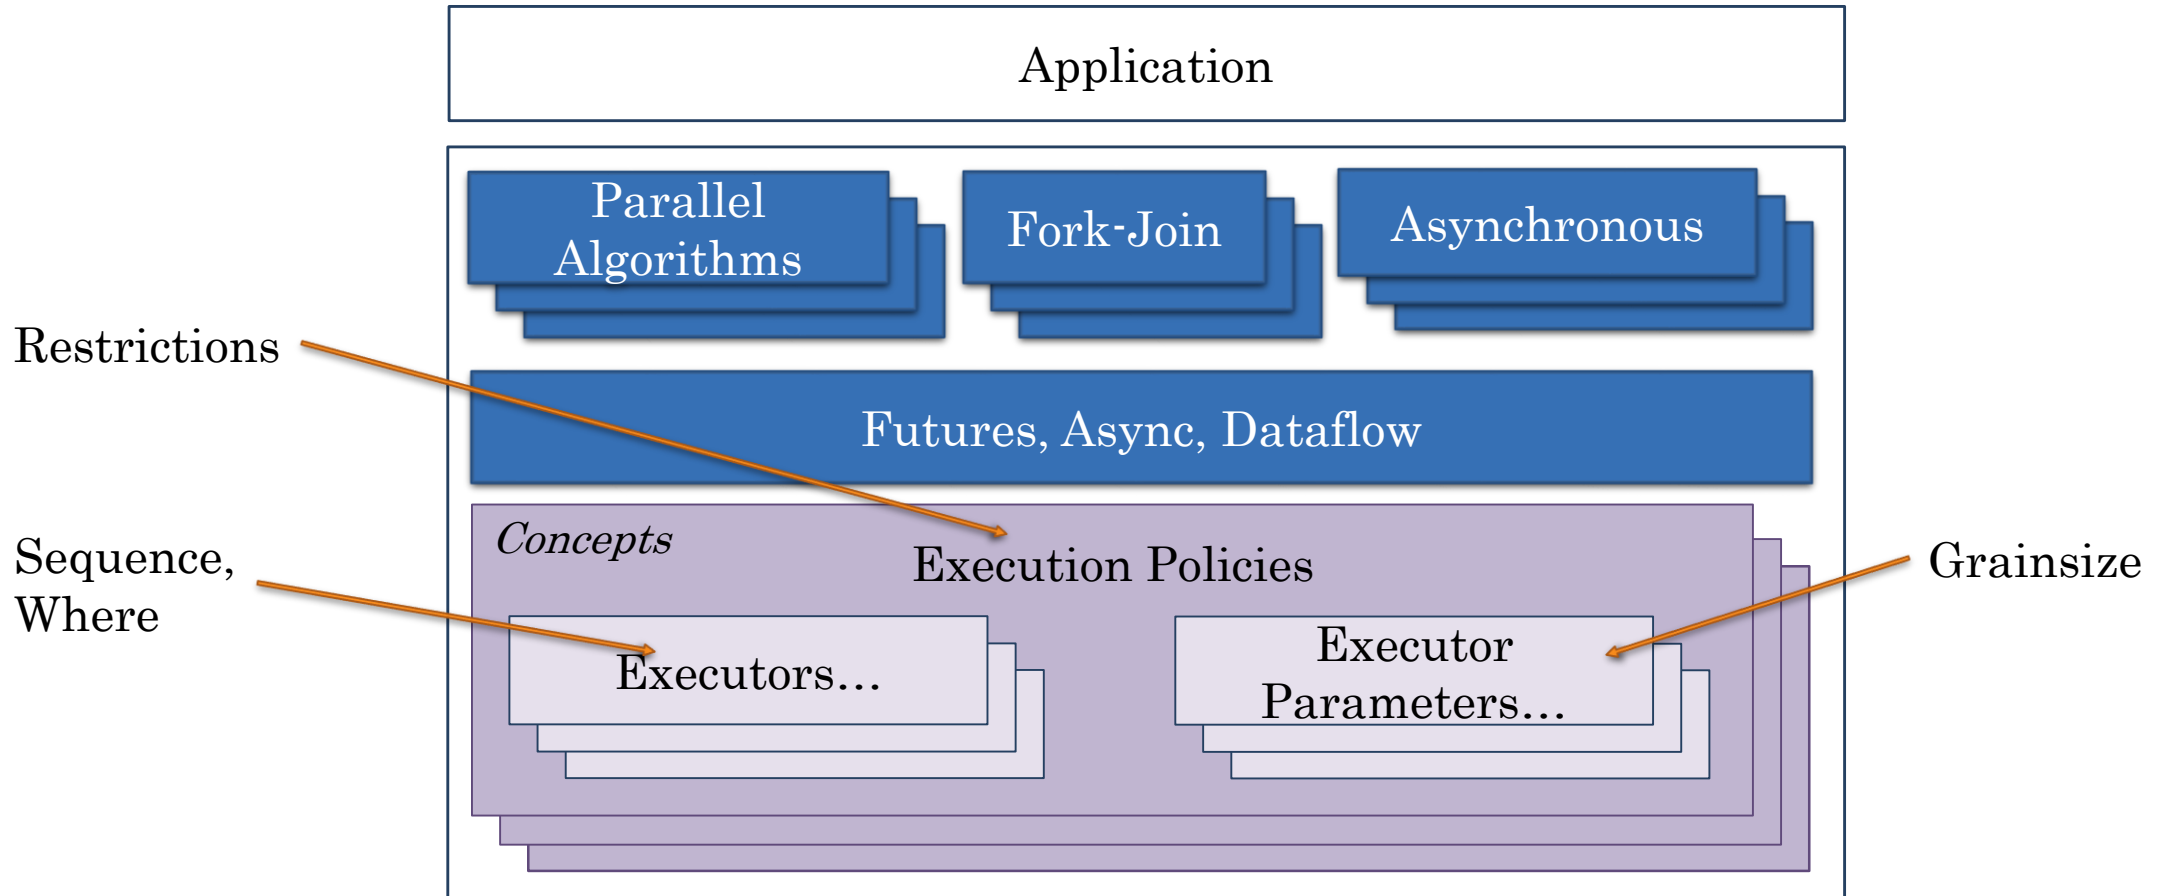
Status of Parallelism in C++

- Current state of standard C++:
 - Parallelism TS: iterative parallelism (moved to be included into C++17)
 - Concurrency TS: task-based, asynchronous, and continuation style parallelism
 - N4411: task blocks for fork-join parallelism of heterogeneous tasks
 - N4406, PR0008R0: executors
 - PR0057R0: resumable functions (co_await, etc.)
- Missing:
 - Integration of the above
 - Parallel ranges
 - Vectorization is being discussed
 - Extensions for GPUs, many-core, distributed, and high-performance computing
- The goal has to be to make parallelism in C++ independent of any external solutions such as OpenMP, OpenACC, etc.
 - HPX makes C++ independent of MPI as well

Controlling What?

- Where and when of execution of code
 - What compute resource to use?
 - In what sequence to execute things?
 - What type of parallelism to apply?
- What restrictions apply to code
 - Can run concurrently, must run sequentially, can be vectorized, etc.
- What additional parameters to apply
 - Grainsize of execution (chunking of loop iterations)

Concepts and Types of Parallelism



HPX

The C++ Standards Library for Concurrency and Parallelism

HPX – A General Purpose Runtime System

- General purpose parallel runtime system for applications of any scale
- Exposes a coherent and uniform, standards-oriented API for ease of programming parallel, distributed, and heterogeneous applications.
 - Enables to write fully asynchronous code using hundreds of millions of threads.
 - Provides unified syntax and semantics for local and remote operations.
- HPX represents an innovative mixture of
 - A global system-wide address space (AGAS - Active Global Address Space)
 - Fine grain parallelism and lightweight synchronization
 - Combined with implicit, work queue based, message driven computation
 - Full semantic equivalence of local and remote execution, and
 - Explicit support for hardware accelerators (through percolation)

HPX – A General Purpose Runtime System

- Enables writing applications which out-perform and out-scale existing applications based on OpenMP/MPI
 - <http://stellar-group.org/libraries/hpx>
 - <https://github.com/STELLAR-GROUP/hpx/>
- Is published under Boost license and has an open, active, and thriving developer community.
- Can be used as a platform for research and experimentation

HPX – The API

- As close as possible to C++11/14/17 standard library, where appropriate, for instance
 - `std::thread` `hpx::thread`
 - `std::mutex` `hpx::mutex`
 - `std::future` `hpx::future` (including N4538, ‘Concurrency TS’)
 - `std::async` `hpx::async` (including N3632)
 - `std::bind` `hpx::bind`
 - `std::function` `hpx::function`
 - `std::tuple` `hpx::tuple`
 - `std::any` `hpx::any` (N3508)
 - `std::cout` `hpx::cout`
 - `std::parallel::for_each`, etc. `hpx::parallel::for_each` (N4507, ‘Parallelism TS’)
 - `std::parallel::task_region` `hpx::parallel::task_region` (N4411)

Parallel Algorithms

Parallel Algorithms

<u>adjacent_difference</u>	adjacent_find	all_of	any_of
copy	copy_if	copy_n	count
count_if	equal	exclusive_scan	fill
fill_n	find	find_end	find_first_of
find_if	find_if_not	for_each	for_each_n
generate	generate_n	includes	inclusive_scan
<u>inner_product</u>	inplace_merge	is_heap	is_heap_until
is_partitioned	is_sorted	is_sorted_until	lexicographical_compare
max_element	merge	min_element	minmax_element
mismatch	move	none_of	nth_element
partial_sort	partial_sort_copy	partition	partition_copy
reduce	remove	remove_copy	remove_copy_if
remove_if	replace	replace_copy	replace_copy_if
replace_if	reverse	reverse_copy	rotate
rotate_copy	search	search_n	set_difference
set_intersection	set_symmetric_difference	set_union	sort
stable_partition	stable_sort	swap_ranges	transform
uninitialized_copy	uninitialized_copy_n	uninitialized_fill	uninitialized_fill_n
unique	unique_copy		

Parallel Algorithms

- Similar to standard library facilities known for years
 - Add execution policy as first argument
- Execution policies have associated default executor and default executor parameters
 - `par` → parallel executor, static chunk size
 - `seq` → sequential executor, no chunking
- Rebind executor and executor parameters:

```
//  
// Simplest case: parallel execution policy  
//  
std::vector<double> d(1000);  
parallel::fill(  
    par,  
    begin(d), end(d), 0.0);
```

Execution Policies (HPX Extensions)

- Extensions: asynchronous execution policies
 - `parallel_task_execution_policy` (asynchronous version of `parallel_execution_policy`), generated with `par(task)`
 - `sequential_task_execution_policy` (asynchronous version of `sequential_execution_policy`), generated with `seq(task)`
- In all cases the formerly synchronous functions return a `future<>`
- Instruct the parallel construct to be executed asynchronously
- Allows integration with asynchronous control flow

Execution Policies (HPX Extensions)

- Extensions: vectorization execution policies
 - `datapar_task_execution_policy` (asynchronous version of `datapar_execution_policy`), generated with `datapar, datapar(task)`
 - `dataseq_task_execution_policy` (asynchronous version of `dataseq_execution_policy`), generated with `dataseq, dataseq(task)`
- Instruct the algorithm to apply certain transformations to used data types allowing for vectorization of code
 - Requires external library: currently Vc (<https://github.com/VcDevel/Vc>), possibly Boost.SIMD
 - Requires use of generic lambdas (C++14) or polymorphic function objects

Executors

Executors

- Executors must implement one function: `async_execute(F && f)`
- Invocation of executors happens through `executor_traits` which exposes (emulates) additional functionality:

```
executor_traits<my_executor_type>::execute(  
    my_executor,  
    [](...){ // perform task },  
    ...);
```

- Four modes of invocation: single async, single sync, bulk async and bulk sync
 - The async calls return a future

Executor Examples

- `sequential_executor`, `parallel_executor`:
 - Default executors corresponding to `par`, `seq`
- `this_thread_executor`
- `distribution_policy_executor`
 - Use one of HPX's (distributed) distribution policies, specify node(s) to run on
- `host::parallel_executor`
 - Specify core(s) to run on (NUMA aware)
- `cuda::default_executor`
 - Use for running things on GPU
- Etc.

Executor Parameters (HPX Extension)

- Same scheme as for executor/executor_traits:
 - parameter/executor_parameter_traits
- Various execution parameters, possibly executor specific
- For instance:
 - Allow to control the grain size of work
 - i.e. amount of iterations of a parallel for_each run on the same thread
 - Similar to OpenMP scheduling policies: static, guided, dynamic
 - auto_chunk_size, static_chunk_size, dynamic_chunk_size
 - Much more fine control
 - Used by parallel algorithms to adjust chunk size
 - Specify GPU-kernel name for certain platforms
 - gpu_kernel<foobar>
 - Specify which other arrays to prefetch

Parallel Algorithms

- Similar to standard library facilities known for years
 - Add execution policy as first argument
- Execution policies have associated default executor and default executor parameters
 - `par` → parallel executor, static chunk size
 - `seq` → sequential executor, no chunking
- Rebind executor and executor parameters:

```
//  
// Simplest case: parallel execution policy  
//  
std::vector<double> d(1000);  
parallel::fill(  
    par,  
    begin(d), end(d), 0.0);
```

Parallel Algorithms

- Similar to standard library facilities known for years
 - Add execution policy as first argument
- Execution policies have associated default executor and default executor parameters
 - `par` → parallel executor, static chunk size
 - `seq` → sequential executor, no chunking
- Rebind executor and executor parameters:

```
// rebind execution policy
//      .on():      executor object, 'where and when'
//      .with():    parameter object(s), possibly executor specific parameters
std::vector<double> d(1000);
parallel::fill(
    par.on(exec).with(par1, par2, ...),
    begin(d), end(d), 0.0);
```

Rebind Execution Policies

```
// uses default execution policy: par
std::vector<double> d = { ... };
parallel::fill(par, begin(d), end(d), 0.0);

// rebind par to user-defined executor
my_executor my_exec = ...;
parallel::fill(par.on(my_exec), begin(d), end(d), 0.0);

// rebind par to user-defined executor and user defined executor parameters
my_params my_par = ...
parallel::fill(par.on(my_exec).with(my_par), begin(d), end(d), 0.0);
```

Data placement

Data Placement

- Different strategies for different platforms
 - Need interface to control explicit placement of data
 - NUMA architectures
 - GPUs
 - Distributed systems
 - Use `std::allocator<T>` interfaces
 - NUMA architectures: first touch
 - Slightly extended: bulk-operations for allocation, construction, destruction, and deallocation

Data Placement

- HPX:
 - `hpx::vector<T, Alloc>`
 - Same interface as `std::vector<T>`
 - Manages data locality through allocator
 - Uses execution target objects for data placement
 - Allows for direct manipulation of data on NUMA domains, GPUs, remote nodes, etc.
 - `hpx::partitioned_vector<T>`
 - Same interface as `std::vector<T>`
 - Segmented data store
 - Segments can be `hpx::vector<T, Alloc>`
 - Uses `distribution_policy` for data placement
 - Allows for manipulation of data on several targets

Data Placement

- Extending `std::allocator_traits`
 - Adding functionality to copy data
 - CPU: trivial
 - GPU: platform specific data transfer, hooked into `parallel::copy`
 - Distributed: maps onto network, possibly RDMA (put/get)
 - Adding functionality to access single elements
 - CPU: trivial
 - GPU: slow, but possible
 - Distributed: maps onto network

Execution Targets

One Ring to Rule them All

Execution Targets

- Opaque types which represent a place in the system
 - Used to identify data placement
 - Used to specify execution site close to data
- Targets encapsulate architecture specifics
 - E.g. `cuda::target`, `host::target`
- Allocators to be initialized from targets
 - Customization of data placement
 - NUMA domain: `host::block_allocator`
 - (possibly remote) GPU device: `cuda::allocator`
 - Locality, i.e. (possibly remote) node
- Executors to be initialized from targets as well
 - Make sure code is executed close to placed data

Examples

Tying it all together

STREAM Benchmark

- Assess memory bandwidth
- Series of parallel for loops, 3 arrays (a, b, c)
 - copy step: $c = a$
 - scale step: $b = k * c$
 - add two arrays: $c = a + b$
 - triad step: $a = b + k * c$
- Best possible performance possible only if data is placed properly
 - Data has to be located in memory of NUMA-domain where thread runs
- OpenMP: implicitly by using ‘first touch’, i.e. run initialization and actual benchmark using same thread
 - `#pragma omp parallel for schedule(static)`

STREAM Benchmark

```
std::vector<double> a, b, c;    // data

// ... init data

auto a_begin = a.begin(), a_end = a.end(), b_begin = b.begin() ...;

// STREAM benchmark
parallel::copy(par, a_begin, a_end, c_begin);           // copy step: c = a
parallel::transform(par, c_begin, c_end, b_begin,       // scale step: b = k * c
    [](double val) { return 3.0 * val; });
parallel::transform(par, a_begin, a_end, b_begin, b_end, c_begin, // add two arrays: c = a + b
    [](double val1, double val2) { return val1 + val2; });
parallel::transform(par, b_begin, b_end, c_begin, c_end, a_begin, // triad step: a = b + k * c
    [](double val1, double val2) { return val1 + 3.0 * val2; });
```


STREAM Benchmark (CPU)

```
host::target tgt("numa=0");           // where and when, here CPU, NUMA domain 0

using executor = host::parallel_executor;
using allocator = host::block_allocator<double>;

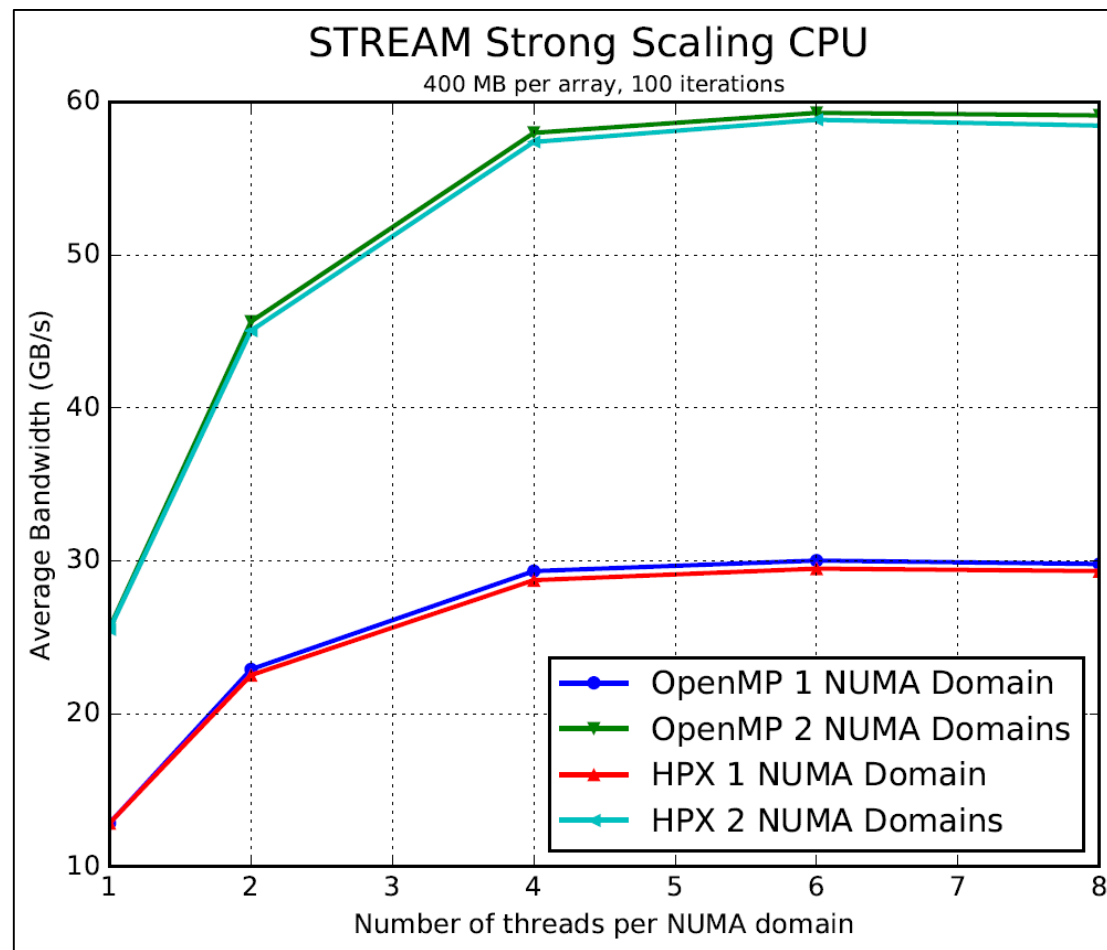
executor exec(tgt);                    // define execution site
allocator alloc(tgt, ...);             // define data placement

vector<double, allocator> a(alloc), b(alloc), c(alloc);           // data
// ... init data

auto policy = par.on(exec).with(static_chunk_size());           // bound execution policy

// STREAM benchmark
parallel::copy(policy, a_begin, a_end, c_begin);
// ...
```

STREAM Benchmark: HPX vs. OpenMP



Extending to GPUs

STREAM Benchmark (GPU)

```
cuda::target target("Tesla C2050"); // where and when, here NVidia GPU (CUDA)

using executor = cuda::default_executor;
using allocator = cuda::allocator<double>;

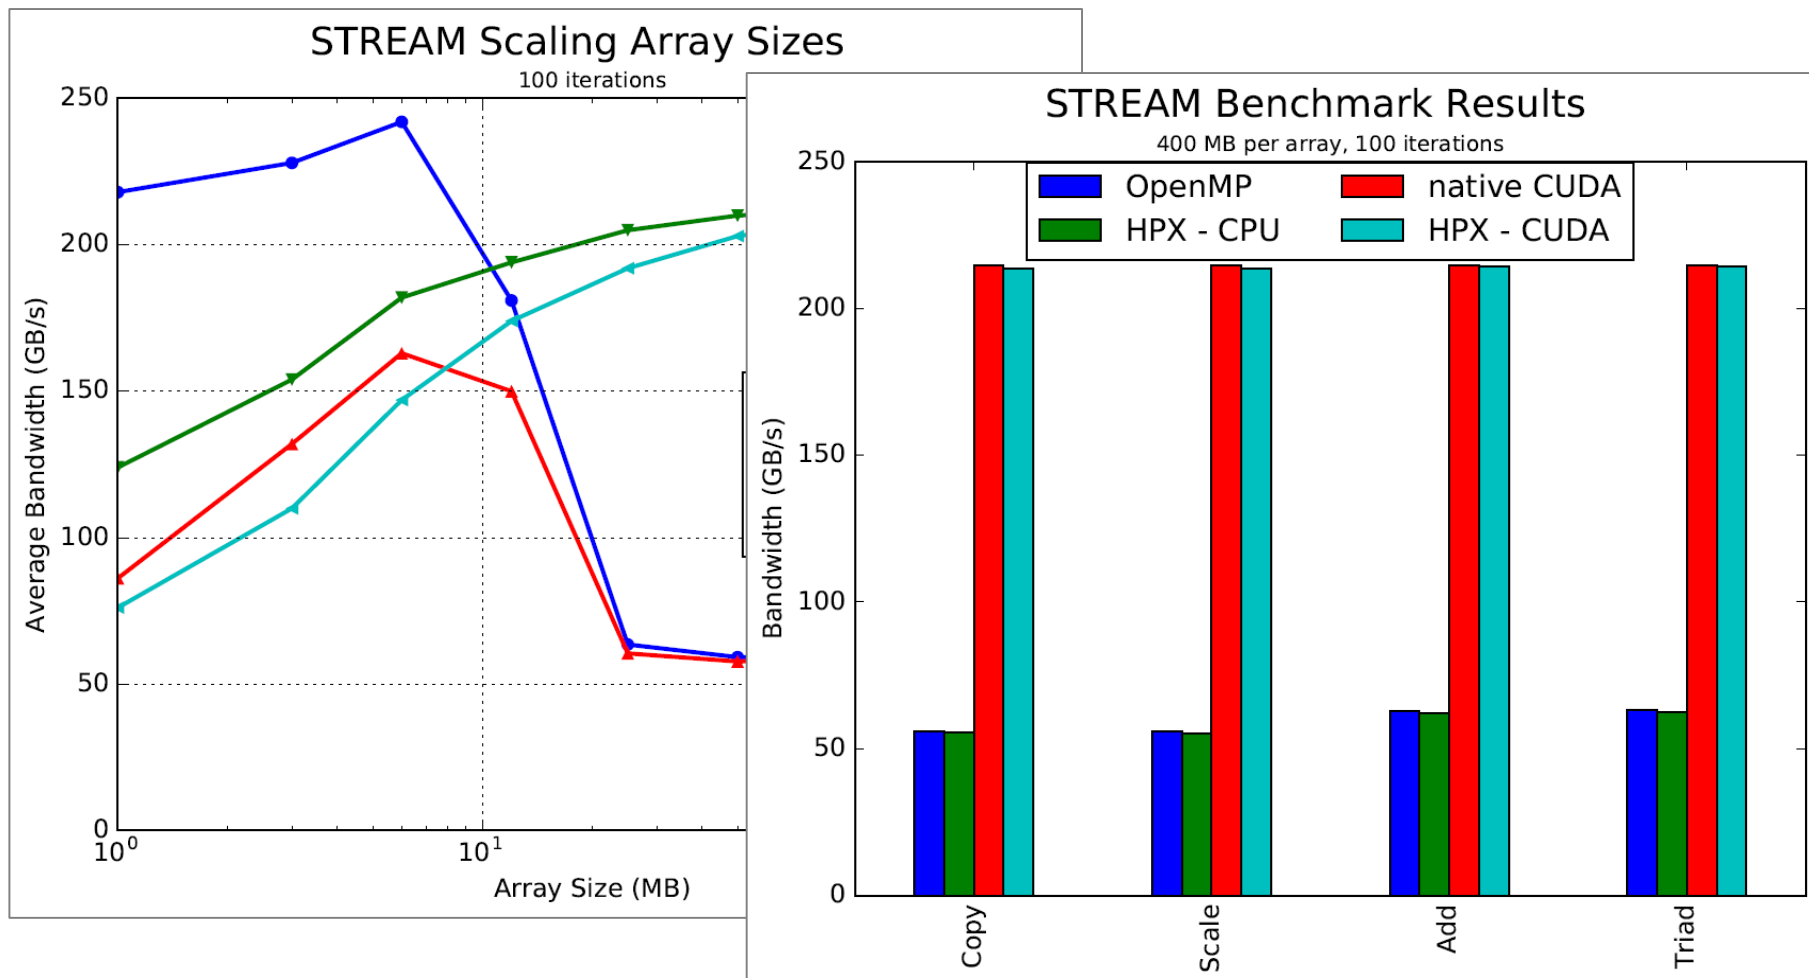
executor exec(tgt);                // define execution site
allocator alloc(tgt);              // define data placement

std::vector<double> data = { ... }; // init data on host
hpx::vector<double, allocator> a(alloc), b(alloc), c(alloc); // data on device

parallel::copy(par, data.begin(), data.end(), a_begin); // copy data to device

// STREAM benchmark
// ...
```

STREAM Benchmark: HPX vs. OpenCL



Vectorization

Dot-product: Parallel Execution

```
std::vector<float> data1 = {...};  
std::vector<float> data2 = {...};  
  
inner_product(  
    par,                                     // just parallel execution  
    std::begin(data1), std::end(data1),  
    std::begin(data2),  
    0.0f,  
    [](auto t1, auto t2) { return t1 + t2; }, // std::plus<>()  
    [](auto t1, auto t2) { return t1 * t2; }  // std::multiplies<>()  
);
```

Dot-product: Vectorization

```
std::vector<float> data1 = {...};  
std::vector<float> data2 = {...};  
  
inner_product(  
    datapar,                                     // parallel and vectorized execution  
    std::begin(data1), std::end(data1),  
    std::begin(data2),  
    0.0f,  
    [](auto t1, auto t2) { return t1 + t2; }, // std::plus<>()  
    [](auto t1, auto t2) { return t1 * t2; }  // std::multiplies<>()  
);
```


Dot-Product: Results

