



**BERKELEY LAB**

LAWRENCE BERKELEY NATIONAL LABORATORY



U.S. DEPARTMENT OF  
**ENERGY**


# The Chombo AMR Framework: Refactoring using AMRStencil as Defensive Programming

Brian Van Straalen



2016 Berkeley C++ Summit

# Block-Structured Adaptive Mesh Refinement (AMR)

- Refined regions are organized into rectangular patches.



- Refinement in time as well as in space for time-dependent problems.
- Local refinement can be applied to any structured-grid data, such as bin-sorted particles.



# Chombo: AMR Software Framework

- **Goal:** to support a wide variety of applications that use AMR by means of a common software framework. Refactoring of **BoxLib** to better enable general usage patterns, such as support for embedded boundary representations of complex geometries.
- **Approach:**
  - Mixed-language programming: C++ for high-level abstractions, Fortran for calculations on rectangular patches.
    - SPMD Distributed memory parallelism, kernels written in Fortran (yuck).
  - Re-useable components, based on mapping of mathematical abstractions to classes. Components are assembled in different ways to implement different applications capabilities.
  - Layered architecture, that hides different levels of detail behind interfaces.
  - Significant effort expended in maintaining professional software development team responsive to a variety of users.
- **Status:** Chombo 3.2 Open Source Release – March, 2013.
  - Chombo 3.3 slated for Q1 2017
  - Currently 1M LOC

# The Problem

- Traditional HPC programming approach: simple programming model.
  - Serial programming on a node
    - vendor / third party compilers. Fortran has been our multidimensional array DSL
  - Bulk-synchronous SPMD programming across nodes
    - MPI Message passing
    - source-level frameworks to hide details of parallelism.
- Simple, long run of success. But, it will not deliver high performance on low-power HPC node architectures of the future **at all scales**.
  - Finer-grain parallelism
    - **library overheads** unacceptable,
    - **heterogeneity** becomes more of a problem.
  - Hierarchical parallelism / co-processing
    - NUMA on a node.
    - Different vendors may require different programming models, applications implementations.
  - **Data movement is more expensive**: both algorithms and software will need to change to perform more floating point operations per byte read / written.
- The programming model to address these concerns has not been worked out yet, but I need to start refactoring Chombo now....

# AMRStencil as a Defensive Programming Model

- AMRStencil: a limited embedded C++11 DSL
  - Main abstractions
    - Stencil, forall, Box, RectMDArray, LevelData
- Abstract away from Chombo Library and Applications:
  - Data placement
  - Dereferencing
  - Iteration
- Allow for multiple parallel execution models to be explored and compared across architectures.

# Multidimensional Arrays

Math:

Tensor-Valued Arrays on Boxes:

$$B \subset \mathbb{Z}^D$$
$$\phi : B \rightarrow \mathbb{T}^C \times \mathbb{T}^D \times \mathbb{T}^E$$

Slices (alias):

$$U : B \rightarrow \mathbb{R}^M$$
$$\text{slice}(U, 1, (q, q)) : B \rightarrow \mathbb{R}$$
$$\text{slice}(U, 1, (q, q))_i = U_i(q)$$

Pointwise operators:

$$\phi : B \rightarrow \mathbb{R}^C \times \mathbb{R}^E$$
$$f : \mathbb{R}^C \times \mathbb{R}^E \rightarrow \mathbb{R}$$
$$f@(\phi) : B \rightarrow \mathbb{R}, f@(\phi)_i = f(\phi_i)$$

C++:

```
RectMDArray<double,CNUM> dirFluxes(ghostBox);  
RectMDArray<double> LaplacianTemp(validBox);  
RectMDArray<double,DIM,DIM> gradu(ghostBox);  
RectMDArray<double,DIM,DIM,DIM> D2vel(validBox);  
RectMDArray<double,DIM> divTau(validBox);
```

```
RectMDArray<double> velcomp = slice<double,DIM,1>  
    (vel, Interval(veldir,veldir));
```

```
forall(fOfPhi,phi,f,B);
```

Array algebra:

```
phi3 = phi1 + phi2; /// etc.
```



# Shifts and Stencils

Shifts:

$$\mathcal{S} = (S_0, \dots, S_{D-1}), S_d(\mathbf{i}) = \mathbf{i} + \mathbf{e}^d$$

Stencil:

$$L = \frac{1}{(\Delta x)^2} \sum_{d=0}^{D-1} \mathcal{S}^{e^d} - 2\mathcal{I} + \mathcal{S}^{-(e^d)}$$

Stencils are not quite first class objects:  
they must be fully specified at compile  
time in order to be efficiently applied  
as operators on arrays (see below).

Stencils have well-defined  
symbolic calculus

$$S_1(S_2(A)) = (S_1 * S_2)(A)$$

$$S_1(c_1 * A) + S_2(A) = (c_1 * S_1 + S_2)(A)$$

C++:

```
array<Shift,DIM> S=getShiftVec();
```

```
Stencil<double> laplace =  
    (-2.*DIM)*(S^zero);  
for (int dir=0;dir<DIM;dir++)  
{  
    Point thishft=getUnitv(dir);  
    laplace = laplace +  
                (S^thishft);  
    laplace = laplace +  
                (S^(thishft*(-1)));  
}  
laplace *= (1.0/a_dx/a_dx);  
}
```

# Applying Stencils to Arrays

Math:

Single-level Stencil:

$$A = L(\phi) \text{ on } B$$
$$A_i = \sum a_s \phi_{i+s}, i \in B$$

Fine to Coarse:

$$\mathcal{P}^{fc*}(U)_i = U_{\lfloor \frac{i}{r} \rfloor}$$
$$\phi^c_+ = \frac{1}{2^D} \mathcal{P}_r^{fc*}(\mathcal{S}^s \phi^f)$$
$$\phi^c_i = \phi^f_{ri+s}$$

Coarse to Fine:

$$\mathcal{P}_r^{cf*}(U)_i = U_{ri}$$
$$\mathcal{S}^s \mathcal{P}_r^{cf*}(\phi^f)_+ = \phi^c$$
$$\phi^f_{ri+s} = \phi^c_i$$

C++:

```
template <class T, int N, int D>
class IplusLaplacian : public Laplacian<T,N,D>
```

```
template <class T, int R>
class restrictSten : public Stencil<T,R,1>
```

```
template <class T, int R>
class prolongSten : public Stencil<T, 1, R>
```

```
RectMDArray<double,1> A = Laplacian<double>(Phi, B);
```

```
PhiC = restrictSten<double, 2>(PhiF, BF);
```

```
PhiF = prolongSten<double, 2>(PhiC, BC);
```



# Pointwise Functions: forall

- Not a new idea. Visit every point in `a_box`

```
template<class T, unsigned int Cdest, unsigned int  
        Csrc, typename Func>
```

```
void forall(RectMDArray<T,Cdest>& a_dest,  
            const RectMDArray<T,Csrc>& a_src,  
            const Func& F, const Box& a_box);
```

- F can be any function that matches the  
`operator[]` signature of `a_dest` and `a_src`
- Did this the first time with `std::bind` a lot, now I use lambdas
- Next version of `RectMDArray` is variadic
  - I don't quite know how to write `forall` with multiple variadic arguments

# Unions of Rectangles (Straight Out of Chombo)

Math:

Union of rectangles (fixed size, defined by bitmap). Syntax for domain boundaries suppressed.

$$\Omega = \bigcup_{i \in \mathcal{B}} \Omega_i$$

$$\Omega_i = [ib_{size}, (i + \mathbf{u})b_{size} - \mathbf{u}] , \mathbf{u} = (1, 1, \dots, 1)$$

Arrays defined over unions of rectangles:

$$\phi_\Omega = \{\phi : \Omega_i \rightarrow \mathbb{R} : i \in \mathcal{B}\}$$

Iterators: assumed to parallel, asynchronous.

Non-blocking communication between data over unions of rectangles.

C++:

```
BoxLayout(unsigned int a_N,  
           vector<Point> a_points);
```

```
BoxLayout bl(s_numblockpower, s_points);  
LevelData<double, 1> phi(bl, s_nghost);
```

```
for(BLIterator blit(layout);  
    blit != blit.end(); ++blit)  
    {initialize(phi[*blit]);}
```

```
phi.exchange();  
phi.copyTo(phi2);
```

# Putting more together: Euler Class excerpts

Tensor type that matches state vector k

inline class member function WToFd

signature for our RK4 integrator

forall Boxes

- local temporary
- apply Stencil
- forall points
  - grab member function with lambda

```

typedef Tensor<double,DIM+2> V;

class Euler{
    inline unsigned int WToFd(V& a_F, const V& a_W,
        int a_d, double gamma);

    void operator()(EulerData& a_k, double a_time, double
        a_dt)
    {
        .
        for(BLIterator BL(a_k.state); BL.ok(); BL++) {
            Box B1 = grow(B0,2);
            RectMDArray<double,DIM+2> W_f(B1);
            W_f |= IplusLaplacian<double, -1, 24>(W, B1);

            forall(F_ave, W_f,
                [this, d, gamma](V& a, const V& b){ return
                WFtoFd(a,b,d,gamma); },
                B_FtoD);
            .
        }
    }
}
    
```

# What is being gained here?

- I can refactor Chombo once. Starting now.
- No parallelism specified
- I can have multiple implementations of AMRStencil
  - serial C++11: written. 4500 LOC
  - UPC++: prototyped
  - MPI distributed memory : prototyped
  - OpenMP non-hierarchical: prototyped (caveat later)
  - Nested OpenMP: prototyped
  - Kokkos: Protonu Basu trying it out
- I can drop instrumentation in at the same time

## 2.2 GHz i7 processor g++ 5.3

-----  
Timer report 0 (42 timers)  
-----

```
[0] root 15.10033 100.0% 1 15691375104 1039.1 MFlops
[1] Euler::advance 14.83863 98.3% 3600 15691276800 1057.5 MFlops
[2] EulerOp::operator 14.17020 93.8% 14400 14544806400 1026.4 MFlops
[3] EulerOp::operator::F_ave 5.14908 34.1% 28800 6559488000 1273.9 MFlops
[4] Stencil::apply 3.43652 22.8% 57600 2156544000 627.5 MFlops
[15] forall_RectMDArray_2 0.90566 6.0% 28800 1437696000 1587.5 MFlops
[22] RectMDArray::plus 0.48670 3.2% 57600 1976832000 4061.7 MFlops
[24] RectMDArray::operator*=:scalar 0.29600 2.0% 57600 988416000 3339.2 MFlops
[5] EulerOp::operator::W_ave_f 1.59491 10.6% 28800 988416000 619.7 MFlops
[6] Stencil::apply 1.59295 10.5% 28800 988416000 620.5 MFlops
[7] EulerOp::operator::W_ave 1.27023 8.4% 14400 1270080000 999.9 MFlops
[12] Stencil::apply 1.10959 7.3% 14400 705600000 635.9 MFlops
[29] RectMDArray::plus 0.15890 1.1% 14400 564480000 3552.4 MFlops
[8] EulerOp::operator::W_f 1.19174 7.9% 28800 718848000 603.2 MFlops
[9] Stencil::apply 1.18962 7.9% 28800 718848000 604.3 MFlops
[10] EulerOp::operator::U 1.11973 7.4% 14400 705600000 630.2 MFlops
[11] Stencil::apply 1.11838 7.4% 14400 705600000 630.9 MFlops
[13] EulerOp::operator::F_bar_f 0.93524 6.2% 28800 1482624000 1585.3 MFlops
[14] forall_RectMDArray_2 0.93301 6.2% 28800 1482624000 1589.1 MFlops
[16] EulerOp::operator::minusDivF 0.88386 5.9% 43200 770457600 871.7 MFlops
```

-----

Timer report 0 (146 timers)

-----

```
[0] root 1.21672 100.0% 1 2337881308 1921.5 MFlops
  [1] Multigrid::vCycle 1.09095 89.7% 9 2173254876 1992.1 MFlops
    [2] Multigrid::pointRelax 0.83451 68.6% 18 1755316224 2103.4 MFlops
      [3] Stencil::apply 0.70748 58.1% 1728 1585446912 2241.0 MFlops
        [10] LevelData::exchange 0.05280 4.3% 108 0 0.0 MFlops
          [12] RectMDArray::copyTo 0.05065 4.2% 22464 0 0.0 MFlops
            [14] RectMDArray::plus 0.02712 2.2% 864 56623104 2087.7 MFlops
              [16] RectMDArray::minus 0.02283 1.9% 864 56623104 2480.4 MFlops
                [17] RectMDArray::operator*=::scalar 0.02236 1.8% 1728 56623104 2532.3MFlops
          [4] Multigrid::vCycle 0.17891 14.7% 9 271662300 1518.4 MFlops
            [5] Multigrid::pointRelax 0.12692 10.4% 18 219414528 1728.7 MFlops
              [6] Stencil::apply 0.10379 8.5% 1728 198180864 1909.4 MFlops
                [20] LevelData::exchange 0.01313 1.1% 108 0 0.0 MFlops
                  [22] RectMDArray::copyTo 0.01122 0.9% 22464 0 0.0 MFlops
                    [32] RectMDArray::plus 0.00473 0.4% 864 7077888 1495.7 MFlops
                      [43] RectMDArray::operator*=::scalar 0.00239 0.2% 1728 7077888 2966.6MFlops
                        [45] RectMDArray::minus 0.00219 0.2% 864 7077888 3228.1 MFlops
                    [13] Multigrid::vCycle 0.03951 3.2% 9 33963228 859.5 MFlops
                      [15] Multigrid::pointRelax 0.02476 2.0% 18 27426816 1107.9 MFlops
```



# How is your sneaky code counting flops?

```
inline unsigned int EulerOp::WToFd(V& a_F, const V& a_W, int a_d, double a_gamma)
    const
{
    double F0 = a_W(a_d+1)*a_W(0);
    double W2 = 0.0;

    a_F(0) = F0;

    for (int d = 1; d <= DIM; d++)
    {
        double Wd = a_W(d);

        a_F(d) = Wd*F0;
        W2 += Wd*Wd;
    }

    a_F(a_d+1) += a_W(DIM+1);

    a_F(DIM+1) = a_gamma/(a_gamma - 1.0) * a_W(a_d+1) * a_W(DIM+1) + 0.5 * F0 * W2;

    return 2*DIM+8;
}
```

← Currently, I'm cheating!

...but this is not a grand cheat


return value can be generated with static analysis

# Summary

- Core execution distilled to agile core functionality in AMRSTencil
- The 1M LOC in Chombo can be programming model/execution model agnostic
  - Separation of Concerns
  - Decent C++ programmers can use AMRStencil
  - Performance Engineers can try new stuff without learning all of Chombo
- BIterator more flexible than it looks
  - order not specified
  - BIterator::Box can be tiled, or threaded
  - Charm++ version of BIterator (Phil Miller)

# Flies in the Ointment

- ...I can *almost* make Stencils constexpr
- Not every Stencil is knowable at compile-time
  - Embedded Boundary Chombo
    - Stencil points and weights from least-squares solve
  - Currently using runtime stencil playback
- MPI and UPC++ init.
  - AMRStencil needs init
- OpenMP thread model
  - BIterator can't launch omp parallel
    - trapped by basic block scope
  - reduction is done with directive, not runtime
  - programming model escaping refactoring



(d) Stencil for a cut cell using weighted least squares.

# What Haven't we Discussed?

- Translation technologies
  - AMRStencil was originally a target for ROSE DSL research. Much care to distinguish compile time vs runtime semantics
  - How much parsing do you really need?
    - full C++11 AST ?
- Tuning
  - Yes, you can implement AMRStencil with meta-programming, but can you tune it?
    - Basu, P., Hall, M., Williams, S., Van Straalen, B., Olier, L., & Colella, P. (2015, May). Compiler-Directed Transformation for Higher-Order Stencils. In *Parallel and Distributed Processing Symposium (IPDPS), 2015 IEEE International* (pp. 313-323). IEEE.